

# Copyright Notice

These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

For the rest of the details of the license, see <https://creativecommons.org/licenses/by-sa/2.0/legalcode>



DeepLearning.AI

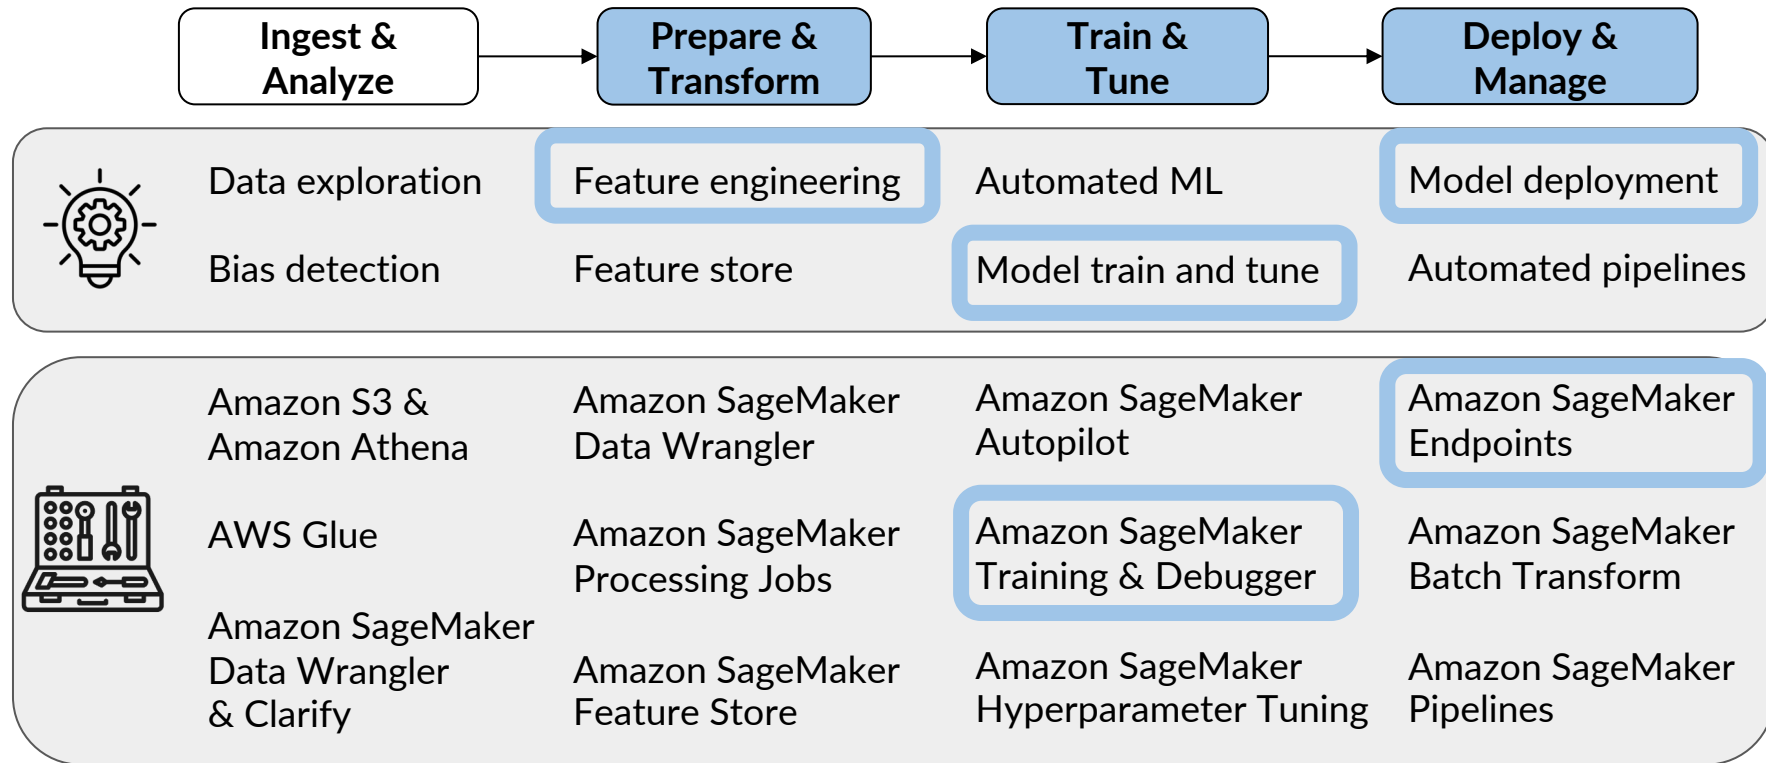


# Practical Data Science

---

## Built-In Algorithms

# Machine Learning Workflow

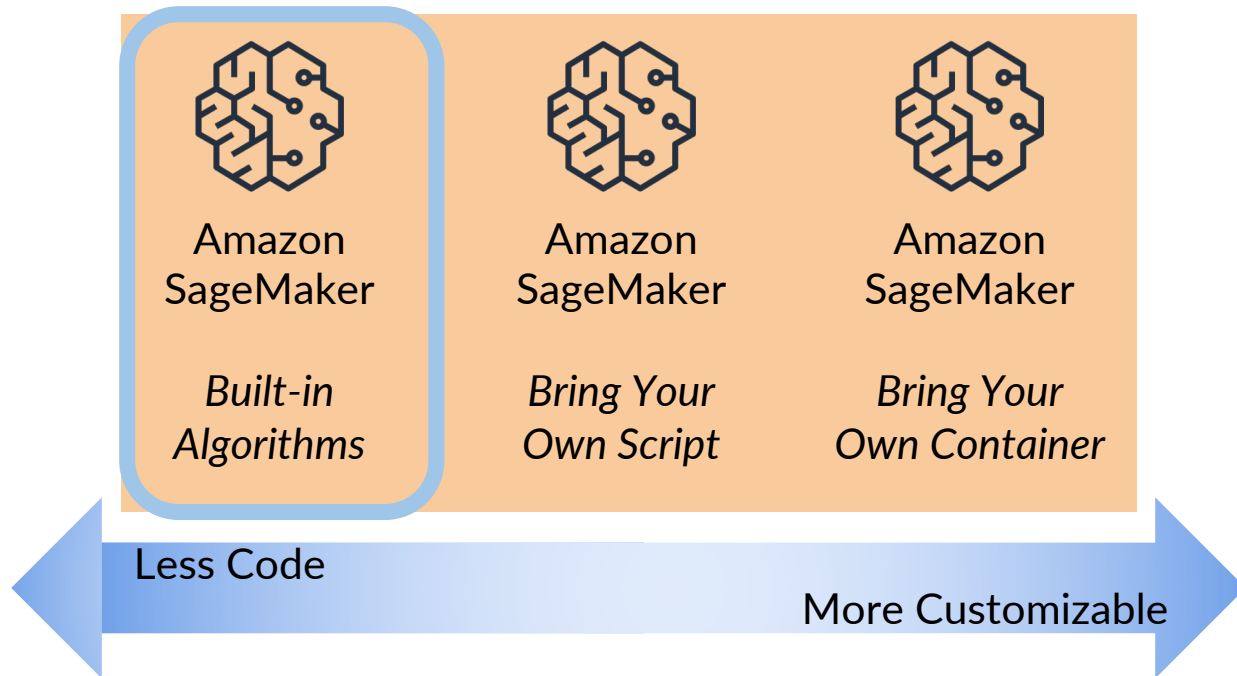


# Why use built-in algorithms?

- Implementations are highly-optimized and scalable
- Focus more on domain-specific tasks rather than managing low-level model code and infrastructure
- Trained model can be downloaded and re-used elsewhere




# When to choose built-in algorithms vs. custom code



# Use cases and algorithms




# Popular ML tasks and learning paradigms



Classification  
& Regression

*Supervised*



Clustering

*Unsupervised*




Image Processing

*Computer Vision*



Text Analysis

*NLP / NLU*

# Classification & regression

Example problems and use cases	Problem types	Input format	Built-in algorithms
Predict if an item belongs to a category: an email spam filter	Binary/multi-class classification	Tabular	XGBoost, K-Nearest Neighbors (k-NN)



# Classification & regression

Example problems and use cases	Problem types	Input format	Built-in algorithms
Predict if an item belongs to a category: an email spam filter	Binary/multi-class classification	Tabular	XGBoost, K-Nearest Neighbors (k-NN)
Predict a numeric/continuous value: estimate the value of a house	Regression	Tabular	Linear Learner, XGBoost

# Classification & regression

Example problems and use cases	Problem types	Input format	Built-in algorithms
Predict if an item belongs to a category: an email spam filter	Binary/multi-class classification	Tabular	XGBoost, K-Nearest Neighbors (k-NN)
Predict a numeric/continuous value: estimate the value of a house	Regression	Tabular	Linear Learner, XGBoost
Predict sales on a new product based on previous sales data	Time-series forecasting	Tabular	DeepAR Forecasting a supervised learning algorithm for forecasting scalar (one dimensional time series) using Recurrent Neural Networks (RNN)

# Clustering

Example problems and use cases	Problem types	Input format	Built-in algorithms
Drop weak features such as the color of a car when predicting its mileage.	Feature engineering: reduce dimensions	Tabular	Principal Component Analysis (PCA)

# Clustering

Example problems and use cases	Problem types	Input format	Built-in algorithms
Drop weak features such as the color of a car when predicting its mileage.	Feature engineering: reduce dimensions	Tabular	Principal Component Analysis (PCA)
Detect abnormal behavior	Anomaly detection	Tabular	Random Cut Forest (RCF)

RCF is unsupervised algorithm for detecting anomalous data points within a dataset.  
RCF associates an anomaly score with each data point. Low score values indicate that the data point is considered normal. High value indicates the presence anomaly in the data.

# Clustering

Example problems and use cases	Problem types	Input format	Built-in algorithms
Drop weak features such as the color of a car when predicting its mileage.	Feature engineering: reduce dimensions	Tabular	Principal Component Analysis (PCA)
Detect abnormal behavior	Anomaly detection	Tabular	Random Cut Forest (RCF)
Group high/medium/low-spending customers from transaction histories	Clustering / grouping	Tabular	K-Means Density Based Clustering algorithm (DBSCAN)

# Clustering

Example problems and use cases	Problem types	Input format	Built-in algorithms
Drop weak features such as the color of a car when predicting its mileage.	Feature engineering: reduce dimensions	Tabular	Principal Component Analysis (PCA)
Detect abnormal behavior	Anomaly detection	Tabular	Random Cut Forest (RCF)
Group high/medium/low-spending customers from transaction histories	Clustering / grouping	Tabular	K-Means
Organize a set of documents into topics based on words and phrases	Topic modeling	Text	Latent Dirichlet Allocation (LDA), Neural Topic Model (NTM)

# Image processing

Example problems and use cases	Problem types	Input format	Built-in algorithms
Content moderation	Image classification	Image	Image Classification

# Image processing

Example problems and use cases	Problem types	Input format	Built-in algorithms
Content moderation	Image classification	Image	Image Classification
Detect people and objects in an image	Object detection	Image	Object Detection



# Image processing

Example problems and use cases	Problem types	Input format	Built-in algorithms
Content moderation	Image classification	Image	Image Classification
Detect people and objects in an image	Object detection	Image	Object Detection
Self-driving cars identify objects in their path	Computer vision	Image	<b>Semantic Segmentation</b> It classifies every pixel in the image. (different from object detection and image classification)

This leads to information such as the shapes of the objects contained in the image. The segmentation output is represented as a grayscale image called a segmentation mask that has the same shape as the input image.

# Text analysis

Example problems and use cases	Problem types	Input format	Built-in algorithms
Convert Spanish to English	Machine translation	Text	Sequence-to-Sequence

# Text analysis

Example problems and use cases	Problem types	Input format	Built-in algorithms
Convert Spanish to English	Machine translation	Text	Sequence-to-Sequence
Summarize a research paper	Text summarization	Text	Sequence-to-Sequence

# Text analysis

Example problems and use cases	Problem types	Input format	Built-in algorithms
Convert Spanish to English	Machine translation	Text	Sequence-to-Sequence
Summarize a research paper	Text summarization	Text	Sequence-to-Sequence
Transcribe call center conversations	Speech-to-text	Text	Sequence-to-Sequence


# Text analysis

Example problems and use cases	Problem types	Input format	Built-in algorithms
Convert Spanish to English	Machine translation	Text	Sequence-to-Sequence
Summarize a research paper	Text summarization	Text	Sequence-to-Sequence
Transcribe call center conversations	Speech-to-text	Text	Sequence-to-Sequence
Classify reviews into categories	Text classification	Text	BlazingText

# Text analysis



# Evolution of text analysis algorithms



Word2Vec  
Jan 2013

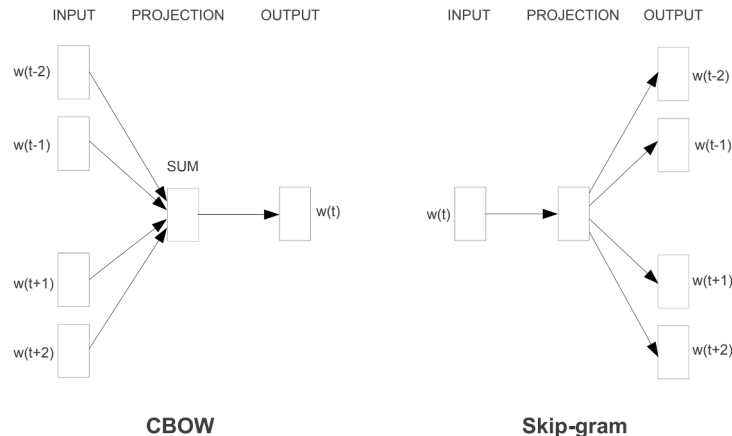
# Text analysis algorithm - Word2Vec

## Concepts

- Convert text into vectors called “embeddings”
- 300-dimensional vector space
- Perform machine learning on the vectors

## Model architectures to create the embeddings

- Continuous bag-of-words (CBOW)
- Continuous skip-gram



Source: "Efficient Estimation of Word Representations in Vector Space", Mikolov et al., 2013



# Evolution of text analysis algorithms

Word2Vec  
Jan 2013

GloVe  
Jan 2014

FastText  
Jul 2016

# Text analysis algorithm - FastText

## Concepts

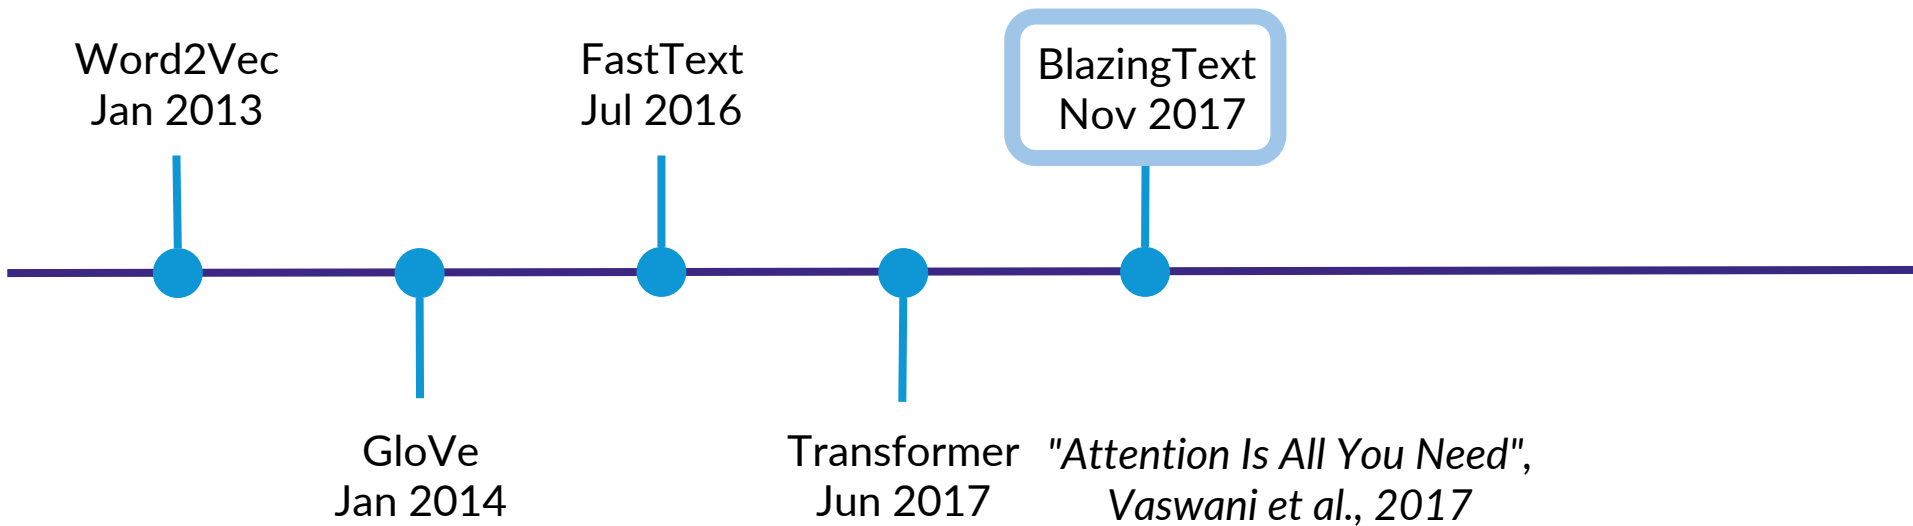
- Extension of word2vec
- Breaks the word into character sets of length n (n-grams):  
"amazon" => "a", "am", "ama", "amaz", "amazo", "amazon"
- Embedding for a word is the aggregate of the embedding of each n-gram within the word

## Implementation

- CBOW and skip-gram models
- Adds text classification


**Helps with the out-of-vocabulary (OOV) issue with word2vec**

# Evolution of text analysis algorithms

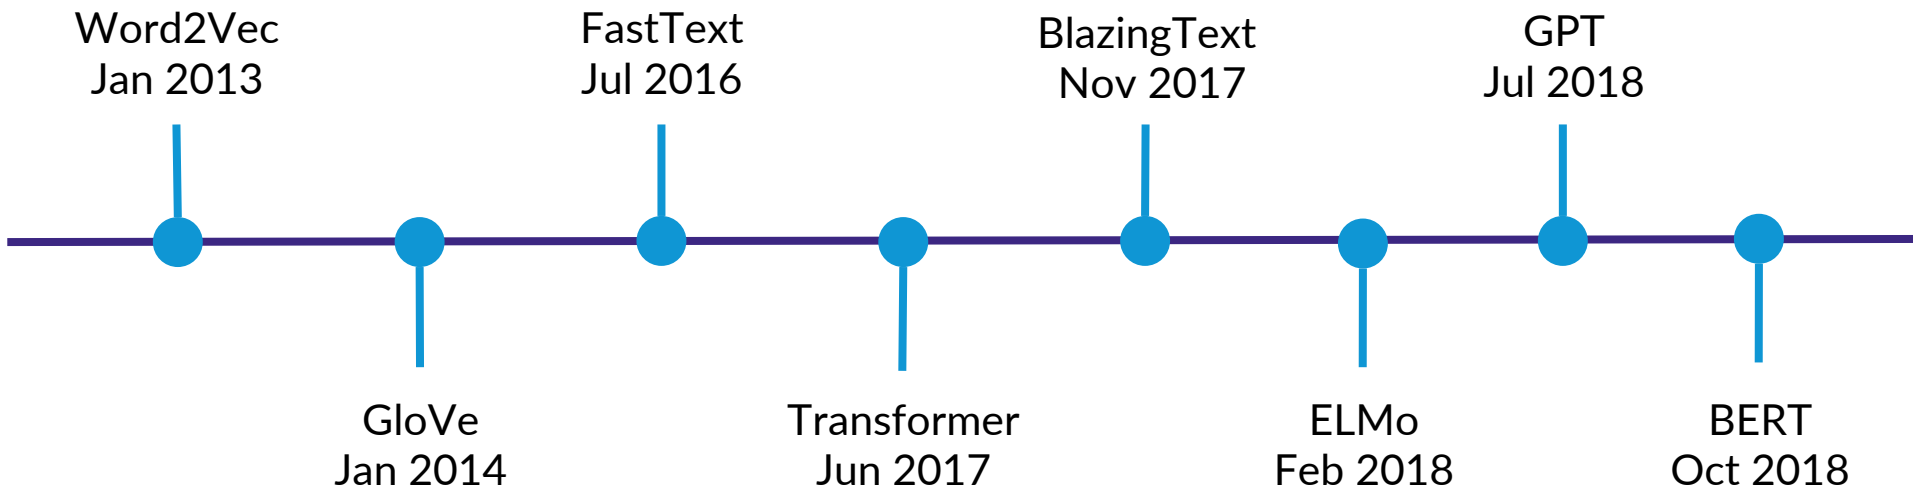


# Text analysis algorithm - BlazingText

- Scales and accelerates Word2Vec using multiple CPUs or GPUs for training
- Extends FastText to use GPU acceleration with custom CUDA kernels
- Creates n-gram embeddings using CBOW and skip-gram
- Saves money by early-stopping a training job
  - when the validation accuracy stops increasing
- Optimized I/O for datasets stored in Amazon S3



# Evolution of text analysis algorithms



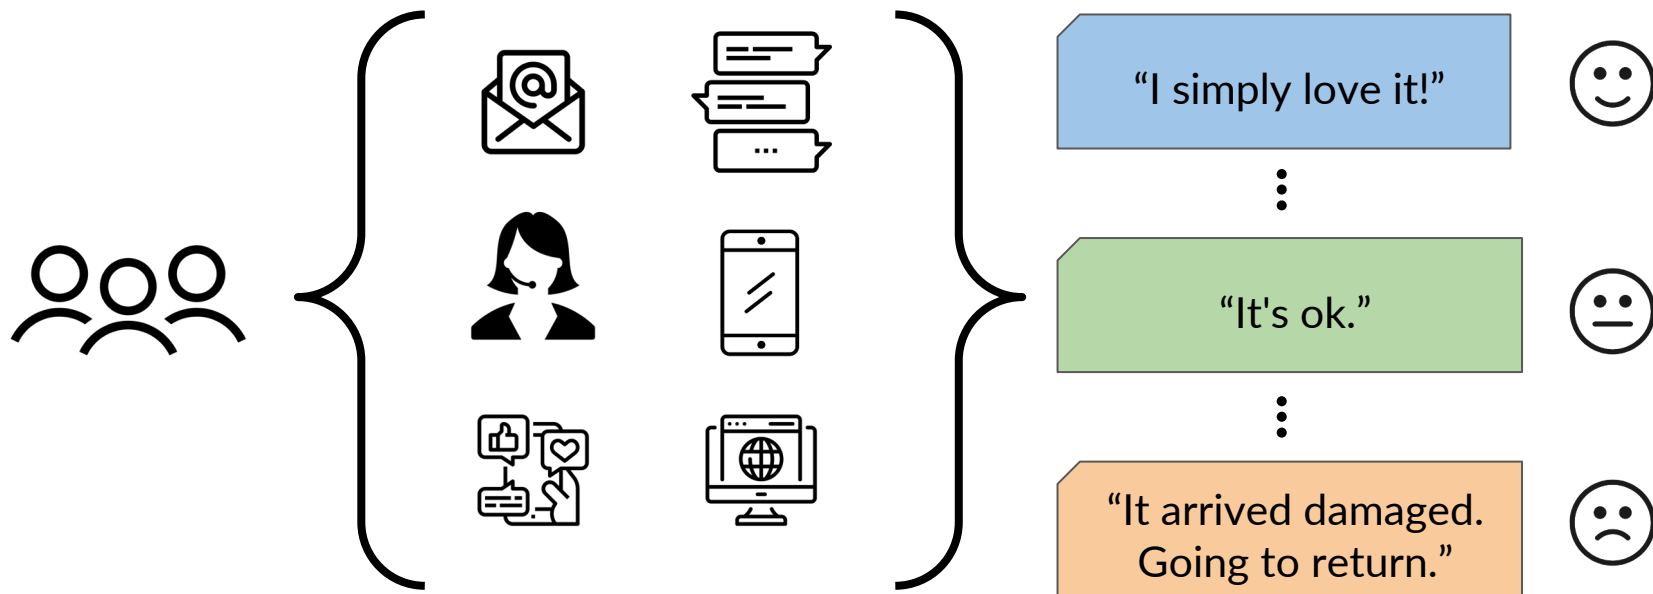
Embeddings from Language models  
Words are learned by a deep bidirectional language model. Hence able to capture syntax and semantics across different linguistic contexts.

# Train a text classifier


with Amazon SageMaker  
BlazingText



# Multi-class classification for sentiment analysis of product reviews




# Multi-class classification for sentiment analysis of product reviews





# Transform raw review data into features



```
sentiment,review_body  
1,"i simply love it"  
0,"it's ok"  
-1,"it arrived damaged. going to return"
```



```
__label__1 "i simply love it ."  
__label__0 "it's ok ."  
__label__-1 "it arrived damaged ."
```



# Transform raw review data into features



NLTK

```
def tokenize(review):  
    return nltk.word_tokenize(review)
```




```
__label__1 "i simply love it ."  
__label__0 "it's ok ."  
__label__-1 "it arrived damaged, going to return ."
```



# Amazon SageMaker BlazingText hyper-parameters for text classification

Parameter Name	Recommended Ranges or Values	Description
epochs	[5-15]	Number of complete passes through the dataset
learning_rate	[0.005-0.01]	Step size for the numerical optimizer
min_count	[0-100]	Discard words that appear less than this number
vector_dim	[32-300]	Number of dimensions in vector space
word_ngrams	[1-3]	Number of words n-gram features to use
early_stopping	True or False	Stop training if validation accuracy stops improving
patience	[5-15]	Number of epochs before early stopping

# Train a text classifier using Amazon SageMaker BlazingText




```
train_data = sagemaker.inputs.TrainingInput(...)
validation_data = sagemaker.inputs.TrainingInput(...)
```

```
data_channels = {
    'train': train_data,
    'validation': validation_data
}
```

```
image_uri = sagemaker.image_uris.retrieve(framework='blazingtext', ...)
```

Retrieves Amazon ECR image URIs for pre-built SageMaker Docker images.

# Train a text classifier using Amazon SageMaker BlazingText



```
train_data = sagemaker.inputs.TrainingInput(...)
validation_data = sagemaker.inputs.TrainingInput(...)
```


```
data_channels = {
    'train': train_data,
    'validation': validation_data
}
```

Retrieves Amazon ECR image URIs for pre-built SageMaker Docker images.

```
image_uri = sagemaker.image_uris.retrieve(framework='blazingtext', ...)
```

```
estimator = sagemaker.estimator.Estimator(image_uri=image_uri, ...)
estimator.set_hyperparameters(...)
estimator.fit(...)
```

# Evaluate the classifier




time	metric_name	value
00.0	train:accuracy	0.4865
10.0	train:accuracy	0.5220
20.0	validation:accuracy	0.5364

# Deploy the text classifier

and make predictions



# Deploy the text classifier




```
text_classifier = estimator.deploy(  
    initial_instance_count=1,  
    instance_type='m1.m4.xlarge', ...)
```

Increase instance\_count > 1 to easily scale out



# Deploy the text classifier



```
text_classifier = estimator.deploy(  
    initial_instance_count=1,  
    instance_type='ml.m4.xlarge', ...)
```

blazingtext-2020-12-07-21-45-06-296

## Endpoint settings

Name	Status
blazingtext-2020-12-07-21-45-06-296	✓ InService
ARN	Creation time
arn:aws:sagemaker:us-east-1:835319576252:endpoint/blazingtext-2020-12-07-21-45-06-296	Mon Dec 07 2020 13:45:07 GMT-0800 (Pacific Standard Time)
	Last updated
	Mon Dec 07 2020 13:51:23 GMT-0800 (Pacific Standard Time)

# Deploy the text classifier



```
text_classifier = estimator.deploy(  
    initial_instance_count=1,  
    instance_type='ml.m4.xlarge', ...)
```

```
payload = {'instances': ['This product is great']}  
response = text_classifier.predict(...)
```

Sample prediction request

# Deploy the text classifier



```
text_classifier = estimator.deploy(  
    initial_instance_count=1,  
    instance_type='ml.m4.xlarge', ...)
```

```
payload = {'instances': ['This product is great']}  
response = text_classifier.predict(...)
```

Sample prediction request

## Sample response:

```
[{  
    "label": ["__label__1"],  
    "prob": [0.9506041407585144]  
}]
```

Prediction response and  
probability score