

# Copyright Notice

These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

For the rest of the details of the license, see <https://creativecommons.org/licenses/by-sa/2.0/legalcode>



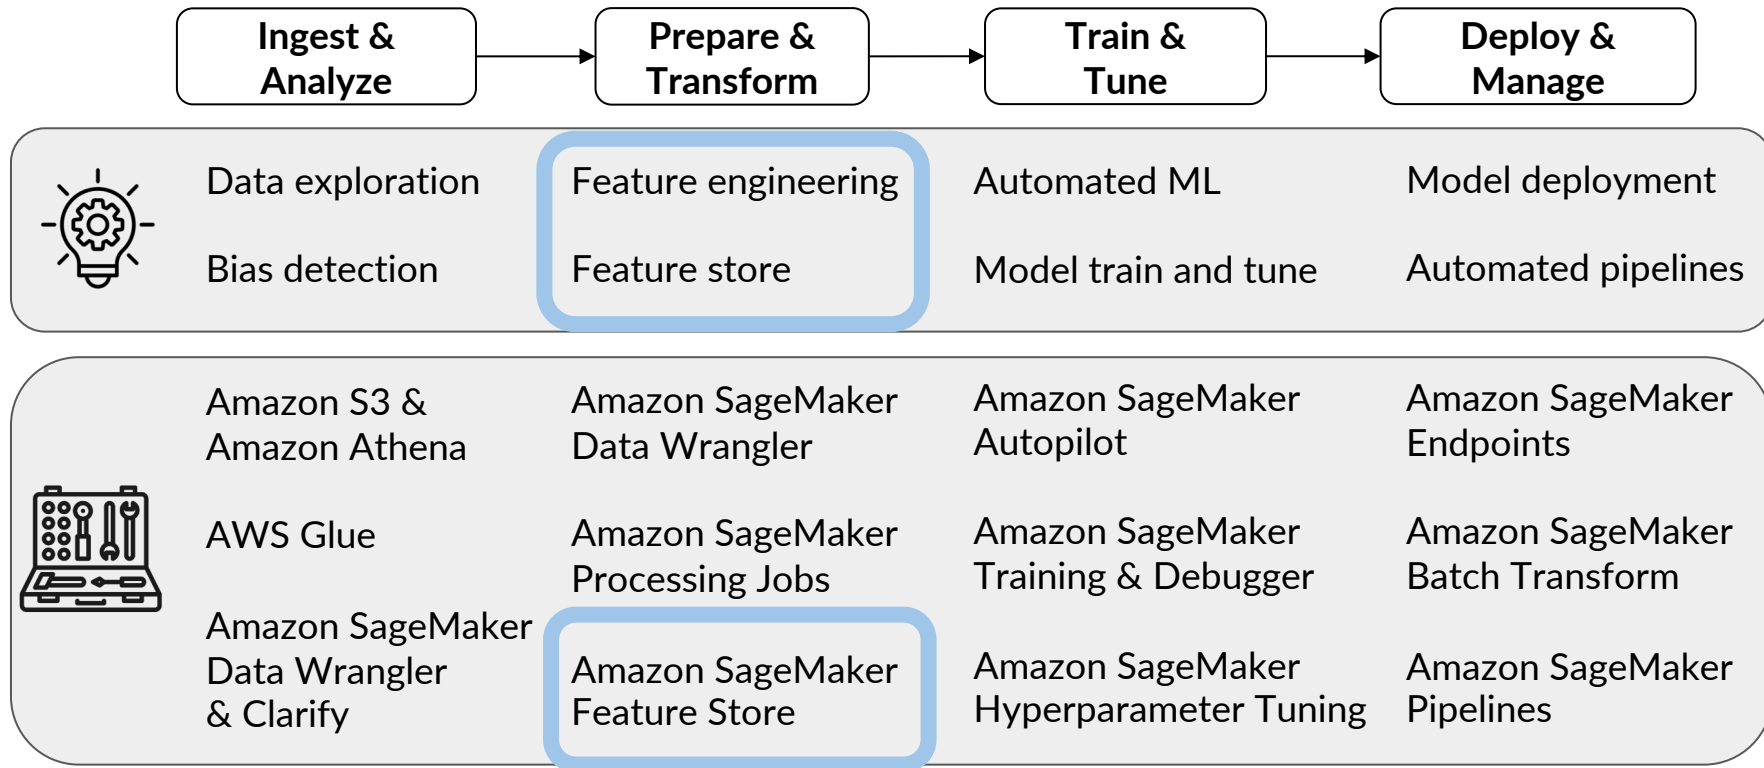
DeepLearning.AI



# Transform Raw Data into Features for Model Training

---

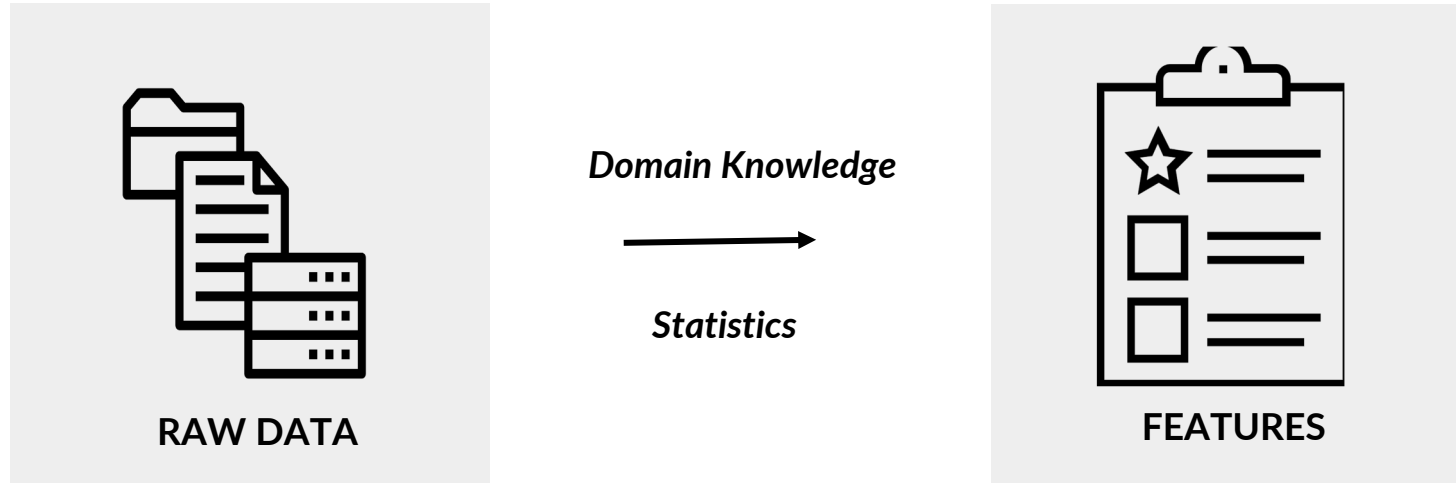
# Machine Learning Workflow



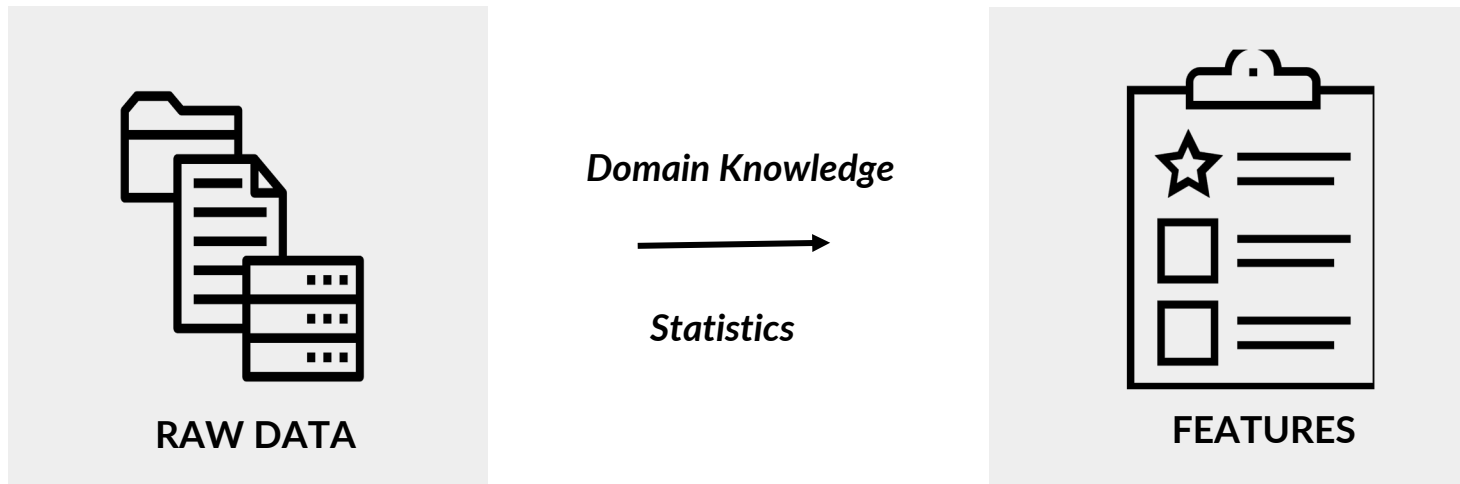
# Feature Engineering



# Feature Engineering




# Feature Engineering



- ✓ Dataset best fits the algorithm
- ✓ Improve ML model performance


# Feature Engineering - Components



Selection





Creation



Transformation

# Feature Engineering - Selection




**Data attributes**

**Irrelevant and redundant attributes**



# Feature Engineering - Selection



**Data attributes**

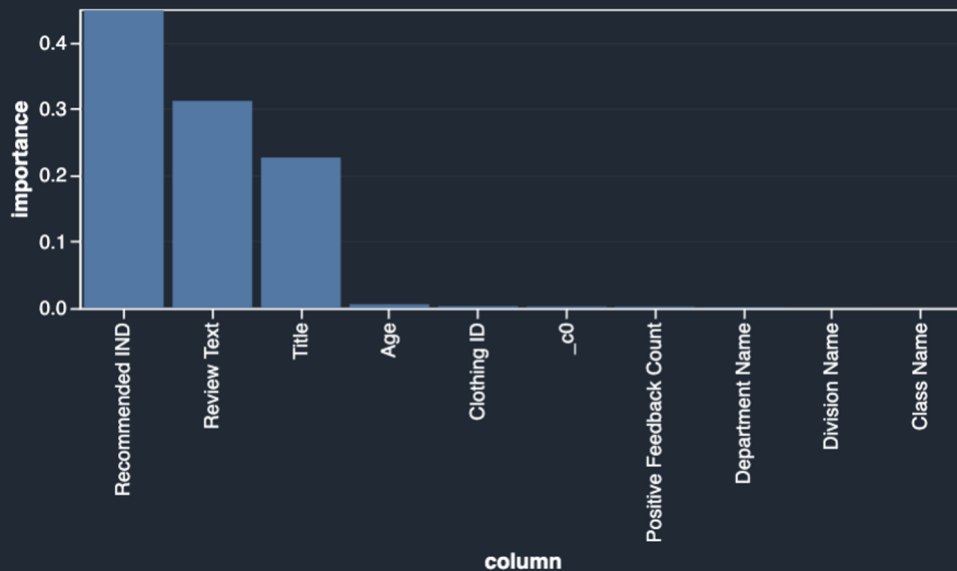
**Irrelevant and redundant attributes**

- ✓ Reduce feature dimensionality
- ✓ Train models faster

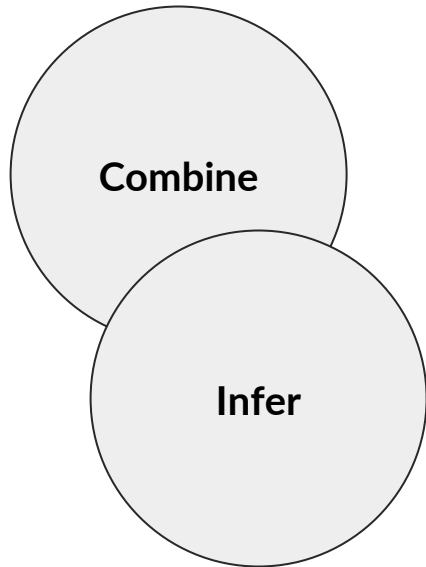

# Feature Importance Report

Quick Model: Product Review Feature Importance

Model achieved a 0.446 f1 on a test set.



# Feature Engineering - Creation




**Existing data points into new features**

**New attributes**




**Lead to more accurate predictions**


# Feature Engineering - Transformation




# Feature Engineering - Transformation



**Missing feature values**




# Feature Engineering - Transformation




**Missing feature values**

**Numerical features**



# Feature Engineering - Transformation



## Missing feature values

Imputation


## Numerical features

Standardization and Normalization


## Non Numerical features

Non Numerical Features are text or category

Categorical Feature can be converted into numeric features by using "one-hot encoding".  
Text need to converted into vectors. more specifically "BERTvectors" or "BERT Embedding"




# Feature Engineering - Transformation



Class Name	Review Text
Blouses	"I simply love it!"
Pants	"It's ok."
Dresses	"It arrived damaged. Going to return."



# Feature Engineering - Transformation



Class Name	Review Text
Blouses	"I simply love it!"
Pants	"It's ok."
Dresses	"It arrived damaged. Going to return."

# Feature Transformation


Review Text

"I simply love it!"
"It's ok."
"It arrived damaged. Going to return."



BERT vectors

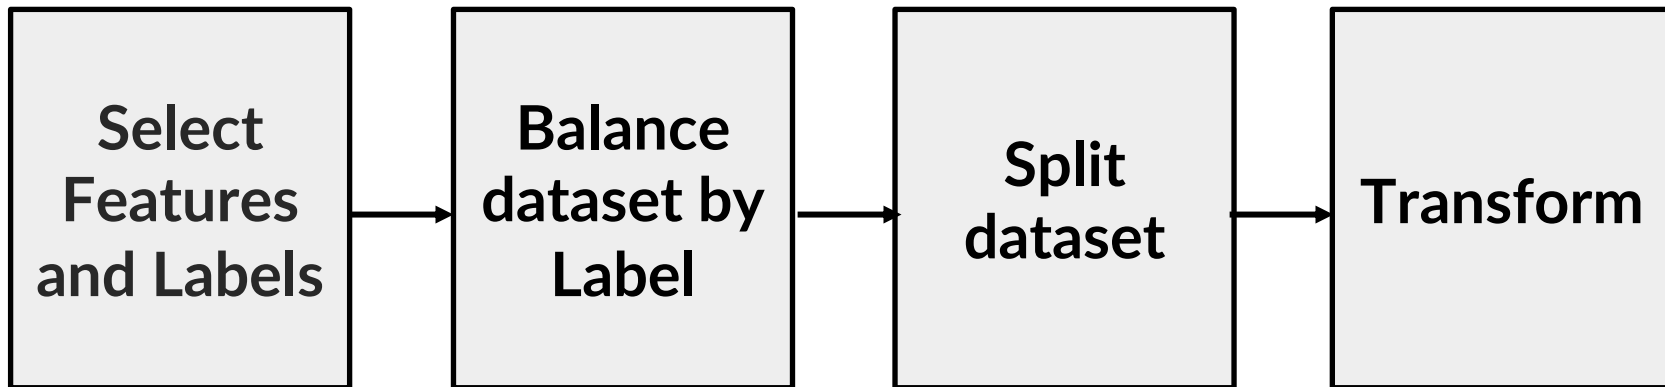
101	2023	...	...
3319	1012	...	...
2003	2307	...	...



# Feature Engineering Pipeline

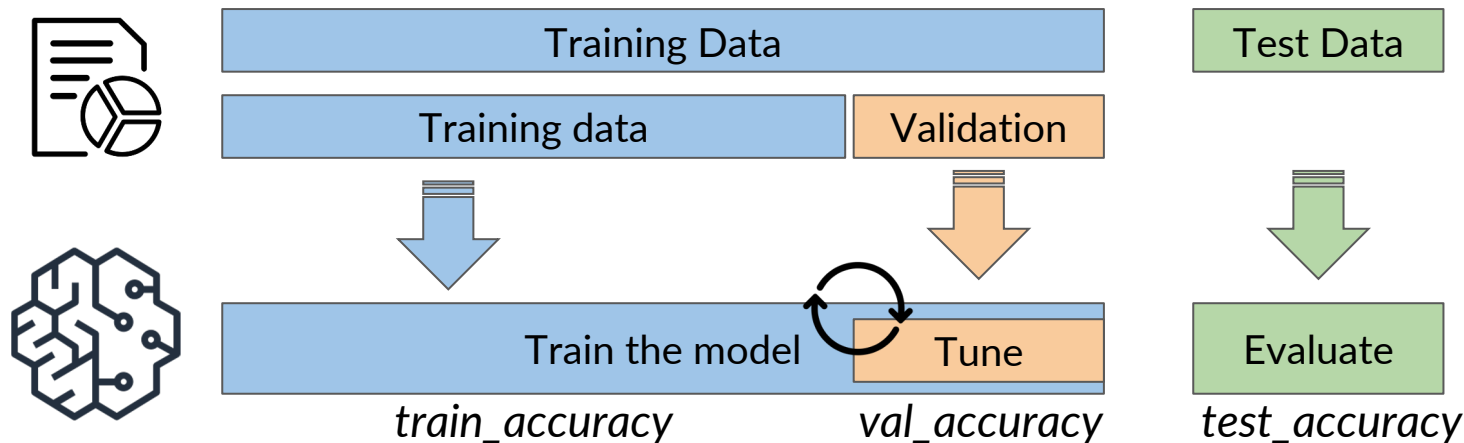


# Feature engineering pipeline

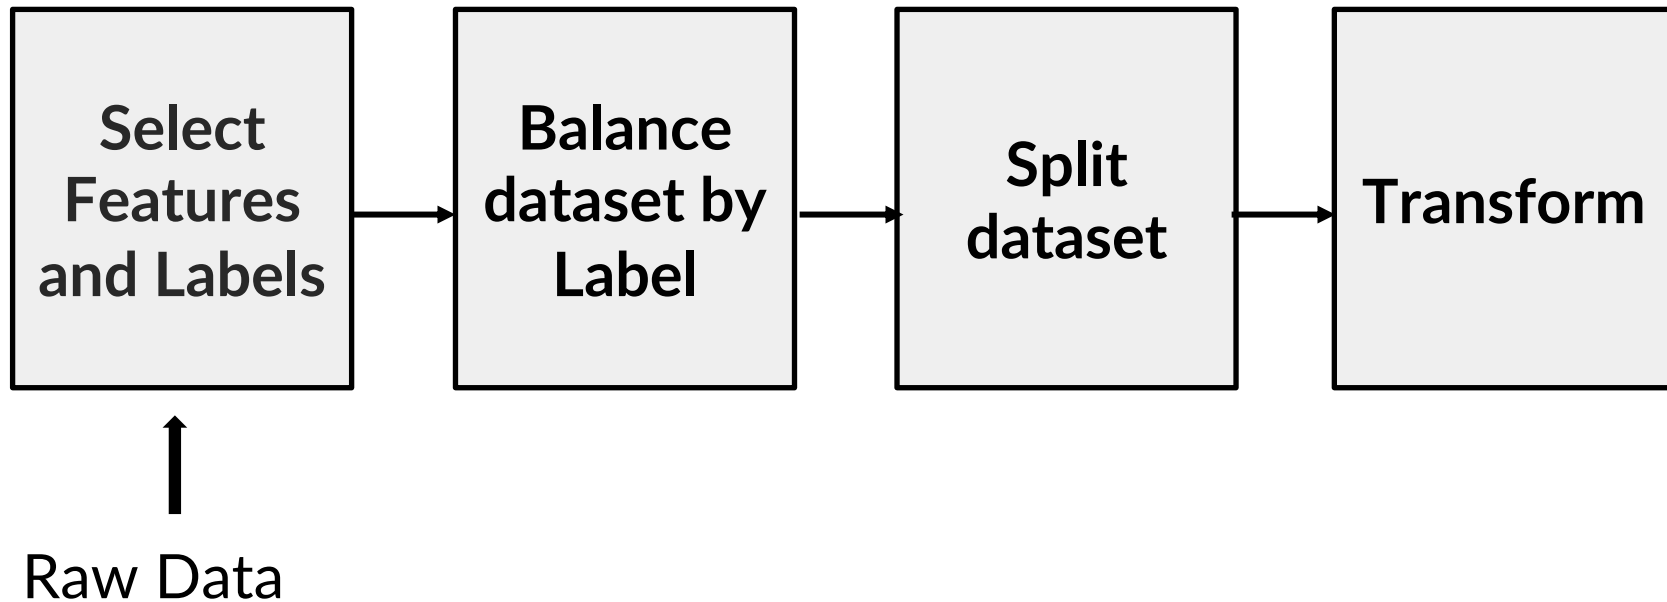


# Split Dataset

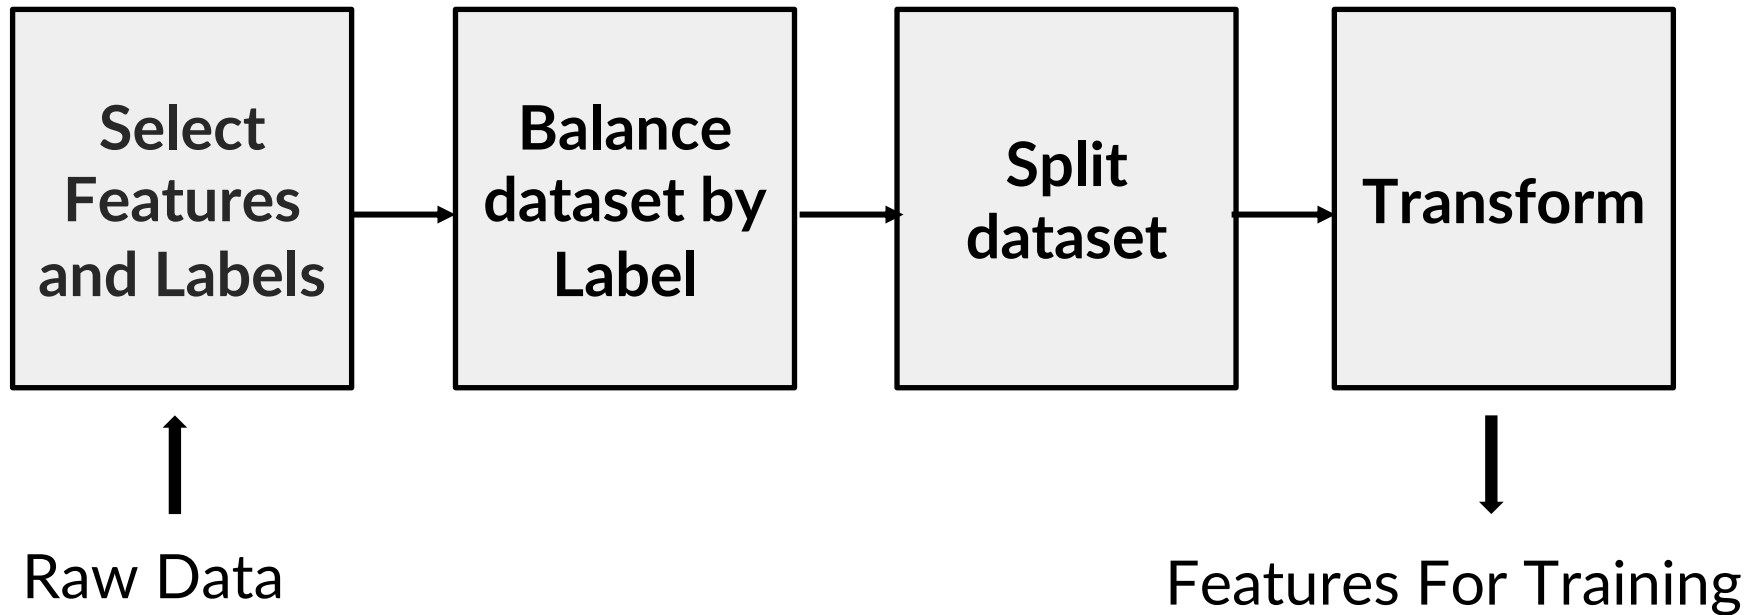
- Training, validation and test data



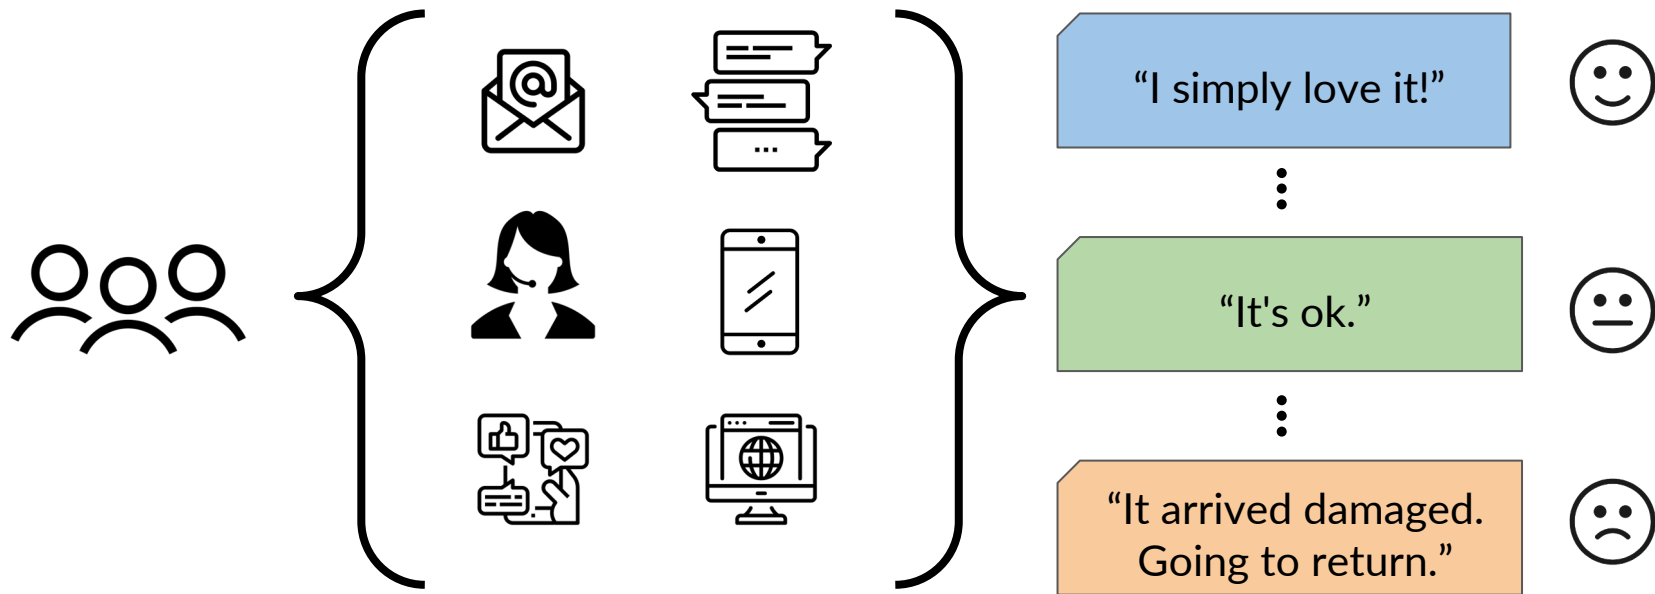
# Feature Engineering Pipeline



# Feature Engineering Pipeline




# Multi-class Classification for Sentiment Analysis of Product Reviews






# Multi-class Classification for Sentiment Analysis of Product Reviews



# Multi-class Classification for Sentiment Analysis of Product Reviews

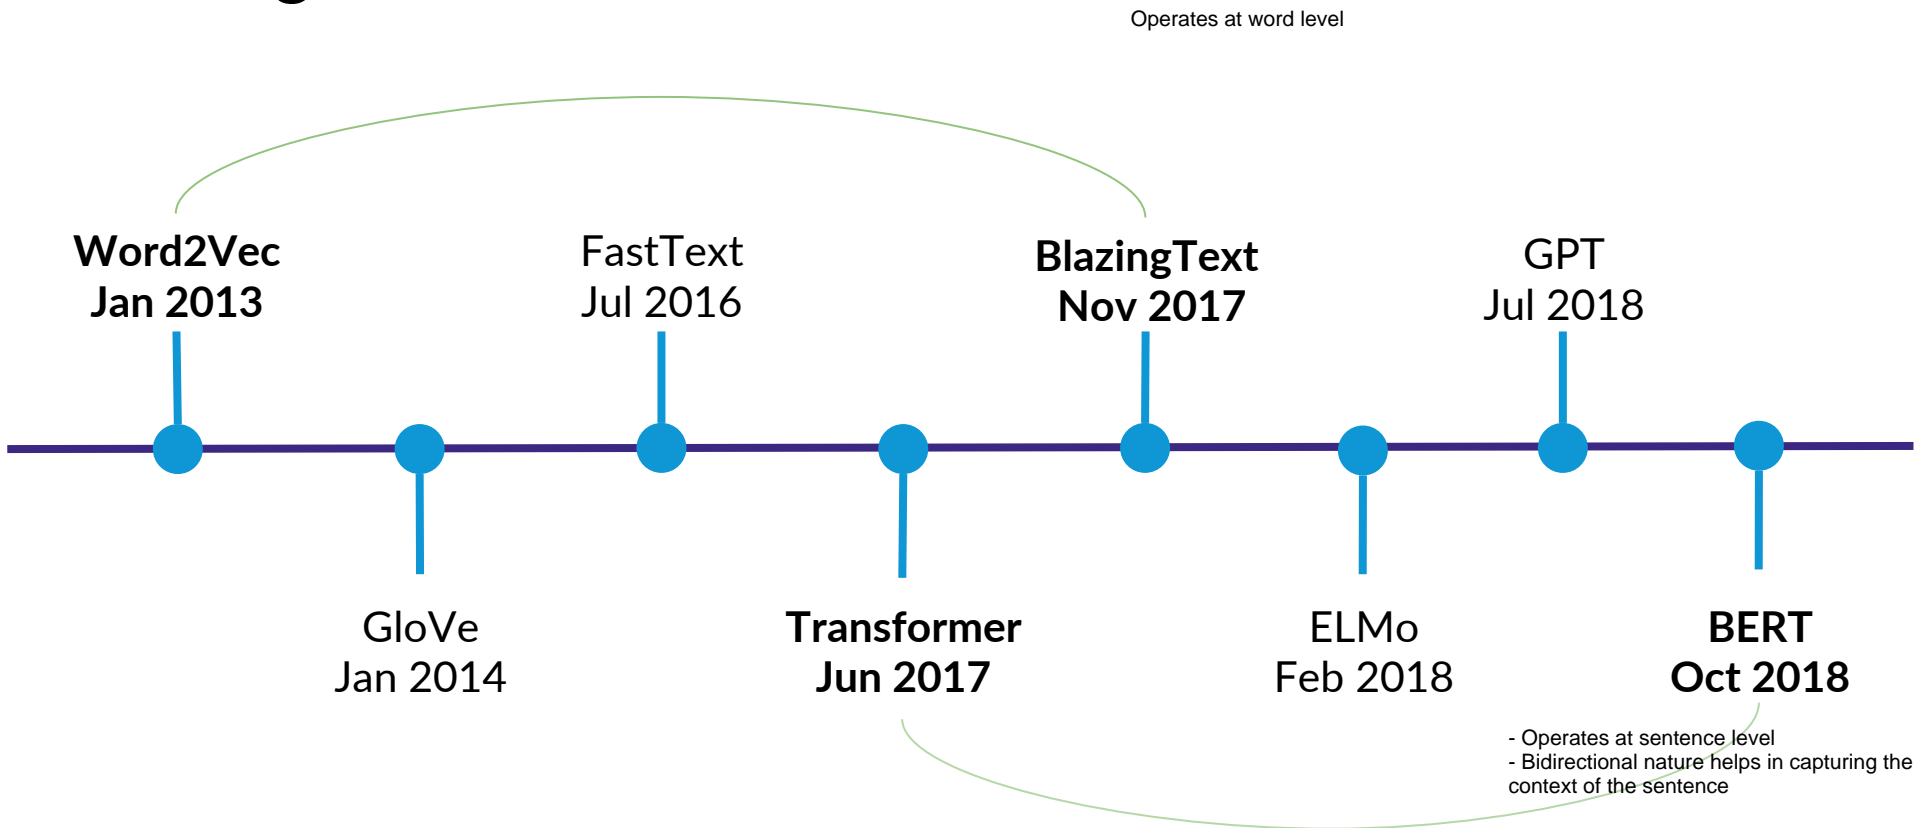


# BERT


Bidirectional Encoder  
Representations  
from Transformers



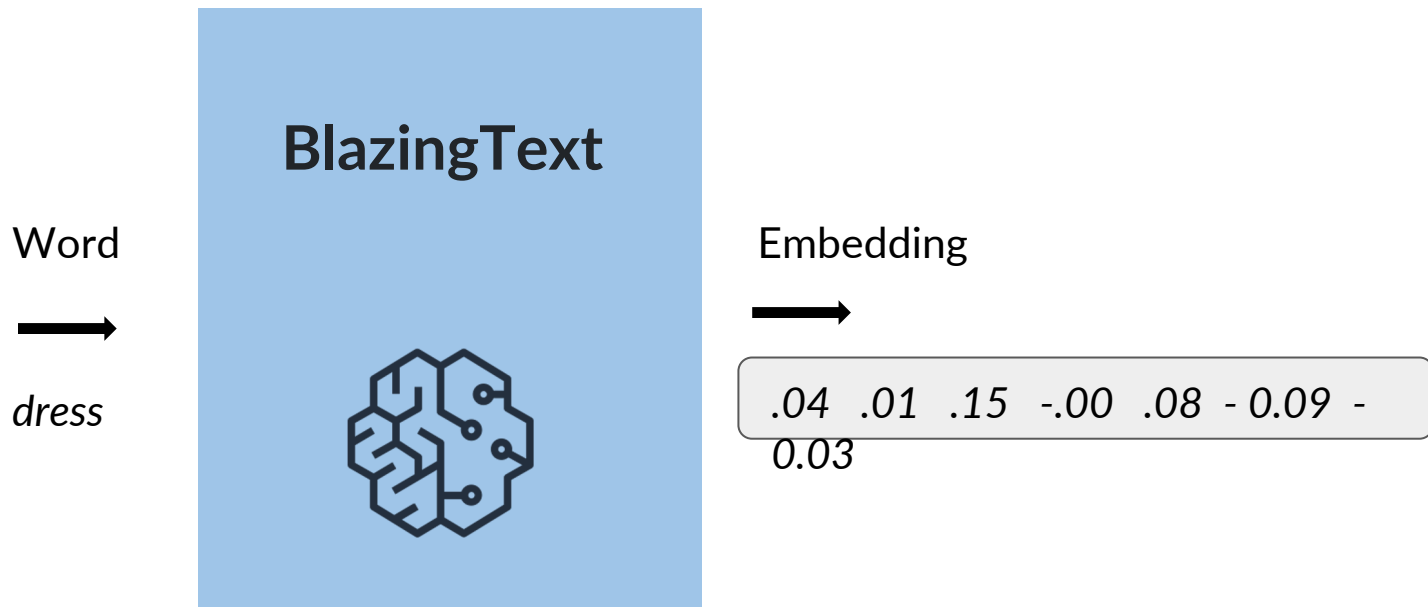
# BlazingText vs BERT




# BlazingText - Word Level Embeddings




# BlazingText - Word Level Embeddings



# BERT - Contextual Embeddings




# BERT - Contextual Embeddings





# BERT - Contextual Embeddings



# BERT Embeddings

Input for BERT Model

(1, 4, 768)

Element wise sum of position,  
segment and token embedding

POSITION EMBEDDING

Input ID



0

1

2

3

(1, 4, 768)

Index position in input sequence

SEGMENT EMBEDDING

Segment ID



0

0

0

0

(1, 4, 768)

0 = Sentence 1  
1 = Sentence 2

TOKEN EMBEDDING

Input ID



101

2293

2023

4377

(1, 4, 768)

Lookup the 768 dimension  
vector dimension

Word Piece  
Tokenization

[CLS], Love, this, dress

1 input sequence  
(consisting of 4 tokens)

Raw Input  
sequence

Love this dress

1 input

# BERT Embeddings

Raw Input  
sequence

Love this dress

**1 input**

# BERT Embeddings

CLS -> indicates Classification problem  
SEP -> token that separates the individual sentences

Word Piece  
Tokenization

(segment words into sub-words with the dimension of 768) + CLS

[CLS], Love, this, dress


**1 input sequence  
(consisting of 4 tokens)**

Raw Input  
sequence


Love this dress

**1 input**


# BERT Embeddings



# BERT Embeddings



# BERT Embeddings



# BERT Embeddings

1 -> 1 input sequence  
4 -> 4 tokens in the sequence  
768 -> 768 dimension vector

Input for BERT Model

(1, 4, 768)

Element wise sum of position,  
segment and token embedding

POSITION EMBEDDING

Input ID



0

1

2

3

(1, 4, 768)

Index position in input sequence

SEGMENT EMBEDDING

Segment ID



0

0

0

0

(1, 4, 768)

0 = Sentence 1  
1 = Sentence 2

TOKEN EMBEDDING

Input ID



101

2293

2023

4377

(1, 4, 768)

Lookup the 768 dimension  
vector dimension

Word Piece  
Tokenization

[CLS], Love, this, dress

1 input sequence  
(consisting of 4 tokens)

Raw Input  
sequence

Love this dress

1 input



# Feature Engineering

At scale with Amazon  
SageMaker Processing Jobs



# RoBERTa model

RoBERTa is built on top of BERT model, but it modifies a few hyper parameters and the way the model is trained.

It also uses a lot more training data than the original BERT model.

## RoBERTa: A Robustly Optimized BERT Pretraining Approach

**Yinhan Liu<sup>\*,§</sup> Myle Ott<sup>\*,§</sup> Naman Goyal<sup>\*,§</sup> Jingfei Du<sup>\*,§</sup> Mandar Joshi<sup>†</sup>  
Danqi Chen<sup>§</sup> Omer Levy<sup>§</sup> Mike Lewis<sup>§</sup> Luke Zettlemoyer<sup>†§</sup> Veselin Stoyanov<sup>§</sup>**

<sup>†</sup> Paul G. Allen School of Computer Science & Engineering,  
University of Washington, Seattle, WA  
{mandar90,lsz}@cs.washington.edu

<sup>§</sup> Facebook AI  
{yinhanliu,myleott,naman,jingfeidu,  
danqi,omerlevy,mikelewis,lsz,ves}@fb.com

### Abstract

Language model pretraining has led to significant performance gains but careful comparison between different approaches is challenging. Training is computationally expensive, often done on private datasets of different

We present a replication study of BERT pre-training (Devlin et al., 2019), which includes a careful evaluation of the effects of hyperparameter tuning and training set size. We find that BERT was significantly undertrained and propose an improved recipe for training BERT models, which

26 Jul 2019

# BERT Embeddings with RoBERTa

```
from transformers import RobertaTokenizer
```

```
PRE_TRAINED_MODEL_NAME = 'roberta-base'
```

Import the  
Tokenizer class

```
tokenizer =  
RobertaTokenizer.from_pretrained(PRE_TRAINED_MODEL_NAME)
```

Create the tokenizer to use  
based on pre trained model

# BERT Embeddings with scikit-learn

encode\_plus method

```
def convert_to_bert_input_ids(...):  
    encode_plus = tokenizer.encode_plus(  
        review,  
        add_special_tokens=True,  
        max_length=128,  
        return_token_type_ids=False,  
        padding='max_length',  
        return_attention_mask=True,  
        return_tensors='pt',  
        truncation=True
```

**Special  
tokens**

Add special tokens or not


**Review to be  
encoded**

**Max sequence  
length**

Defines the max length sequence.

```
    return encode_plus['input_ids'].flatten().tolist()
```

# BERT hyper-parameter: max\_seq\_length



mean	52.51
std	31.38
min	1.00
10%	10.00
20%	22.00
30%	32.00
40%	41.00
50%	51.00
60%	61.00
70%	73.00
80%	88.00
90%	97.00

<b>100%</b>	<b>115.00</b>
-------------	---------------

# BERT Embeddings with scikit-learn

```
def convert_to_bert_input_ids(...):
```

```
    encode_plus = tokenizer.encode_plus(
```

```
        review,
```

Review to be  
encoded

```
        add_special_tokens=True,
```

```
        max_length=128,
```

Max sequence  
length

Special  
tokens

```
        return_token_type_ids=False,
```

```
        padding='max_length',
```

```
        return_attention_mask=True,
```

```
        return_tensors='pt',
```

```
        truncation=True
```

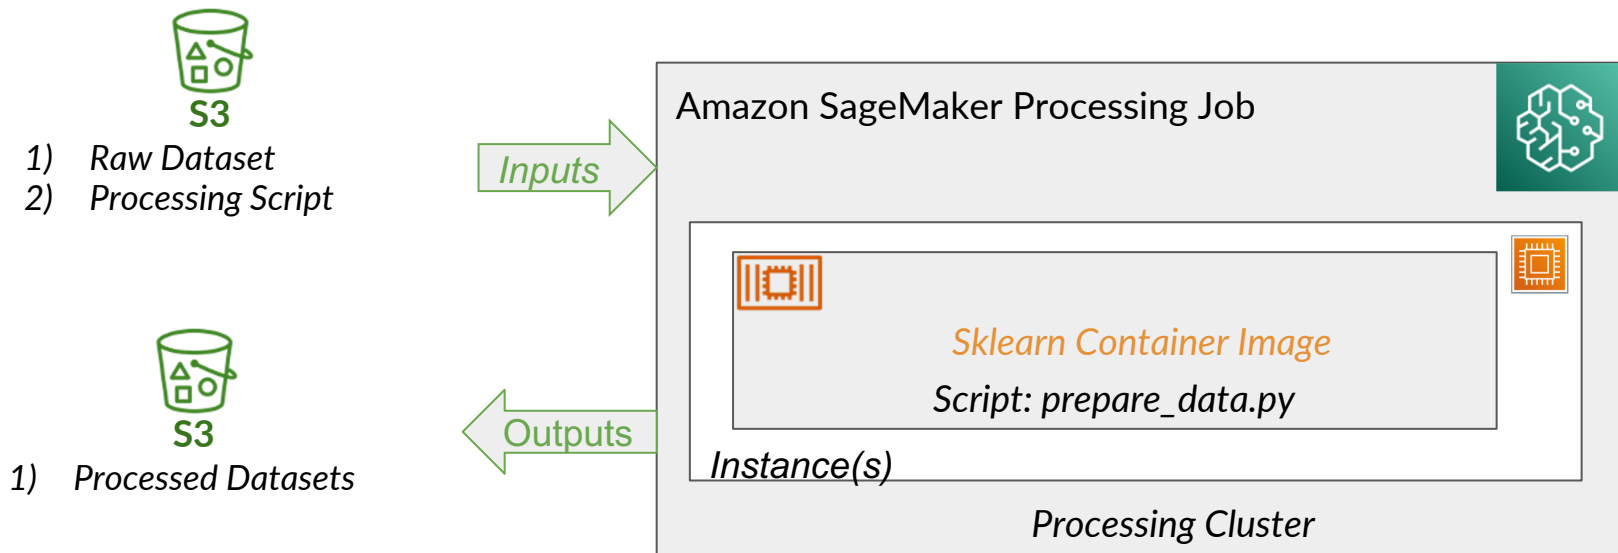
```
    return encode_plus['input_ids'].flatten().tolist()
```

# Amazon SageMaker Processing

Allows us to perform feature engineering at scale

Allows us to perform preprocessing, post-processing and Model evaluation at scale by using a distributed cluster.

## Execute preprocessing, post processing, model evaluation



You can define how many nodes and the types of nodes that you want to include in a cluster.

# Amazon SageMaker Processing with scikit-learn

```
from sagemaker.sklearn.processing import SKLearnProcessor
from sagemaker.processing import ProcessingInput, ProcessingOutput
```

```
processor = SKLearnProcessor(
    framework_version='<SCIKIT_LEARN_VERSION',
    role=role,
    instance_type='ml.c5.4xlarge',
    instance_count=2)
```

**Setup processing  
cluster**

```
processor.run(<parameters>)
```

**Run the  
processing job**



# Amazon SageMaker Processing with scikit-learn

```
...  
code='preprocess-scikit-text-to-bert.py',  
  
inputs=[  
    ProcessingInput(  
        input_name='raw-input-data',  
        source=raw_input_data_s3_uri,  
        ...)  
],
```

**Scikit-learn  
script to execute**

**Input data  
to transform**

# Amazon SageMaker Processing with scikit-learn

```
...
outputs=[
    ProcessingOutput(
        output_name='bert-train',
        s3_upload_mode='EndOfJob',
        source='/opt/ml/processing/output/bert/train'),
    ...,
],
```

Output from the  
processing job

# Amazon SageMaker Processing with scikit-learn

Sentiment	Review
1	<i>this is a great item!</i>
-1	<i>not a good product.</i>
0	<i>dress is ok</i>
-1	<i>do not use! awful. blah</i>



SageMaker  
Processing


label_id	input_ids		
1	101	2023	...
-1	3319	1012	...
0	2003	2307	...
-1	102	3212	...

# Feature Store

Store the results of Feature engineering efforts and reuse those results, so you don't have to run the feature engineering pipeline again and again.




# Feature Store




Centralized

# Feature Store



Centralized




Reusable


# Feature Store

Multiple teams can contribute their features to this centralized repository

Reuse of engineered features, not just across multiple phases of a single machine learning project, but across multiple learning projects.




Centralized




Reusable

Any team member can come in and search for the features they want, and use the search results in their own machine learning projects.



Discoverable

# Feature Engineering Pipeline Extended






# Amazon SageMaker Feature Store




# Amazon SageMaker Feature Store




Store and Serve  
Features

# Amazon SageMaker Feature Store




Store and Serve  
Features




Reduce skew


# Amazon SageMaker Feature Store



Store and Serve  
Features



Reduce skew



Real time & Batch

# Amazon SageMaker Feature Store - Create

```
from sagemaker.feature_store.feature_group import FeatureGroup

reviews_feature_group_name = "reviews_distilbert_max_seq_length_128"

reviews_feature_group = FeatureGroup(
    name=...,
    feature_definitions=...,
    sagemaker_session=sagemaker_session)

reviews_feature_group.create(
    s3_uri="s3://{}/{}".format(bucket, prefix),
    record_identifier_name=record_identifier_feature_name,
    event_time_feature_name=event_time_feature_name,
    role_arn=role)
```

Name

Create

Feature Group is a construct that allows you to group multiple features together and treat them as a set.

First, you define a feature group.Name, definitions, and sagemaker session definition -> name and type of the features.

Create method expects an s3 location where the feature group, along with the individual features will be saved.

# Amazon SageMaker Feature Store - Ingest

```
reviews_feature_group.ingest(  
    data_frame=df_records,  
    max_workers=3,  
    wait=True)
```



**Ingest**

Ingest API is used to ingest feature into the feature group in a multi-threaded fashion

# Amazon SageMaker Feature Store - Retrieve

```
reviews_feature_store_query =  
    reviews_feature_group.athena_query()
```

**Query S3**

```
reviews_feature_store_table =  
    reviews_feature_store_query.table_name
```

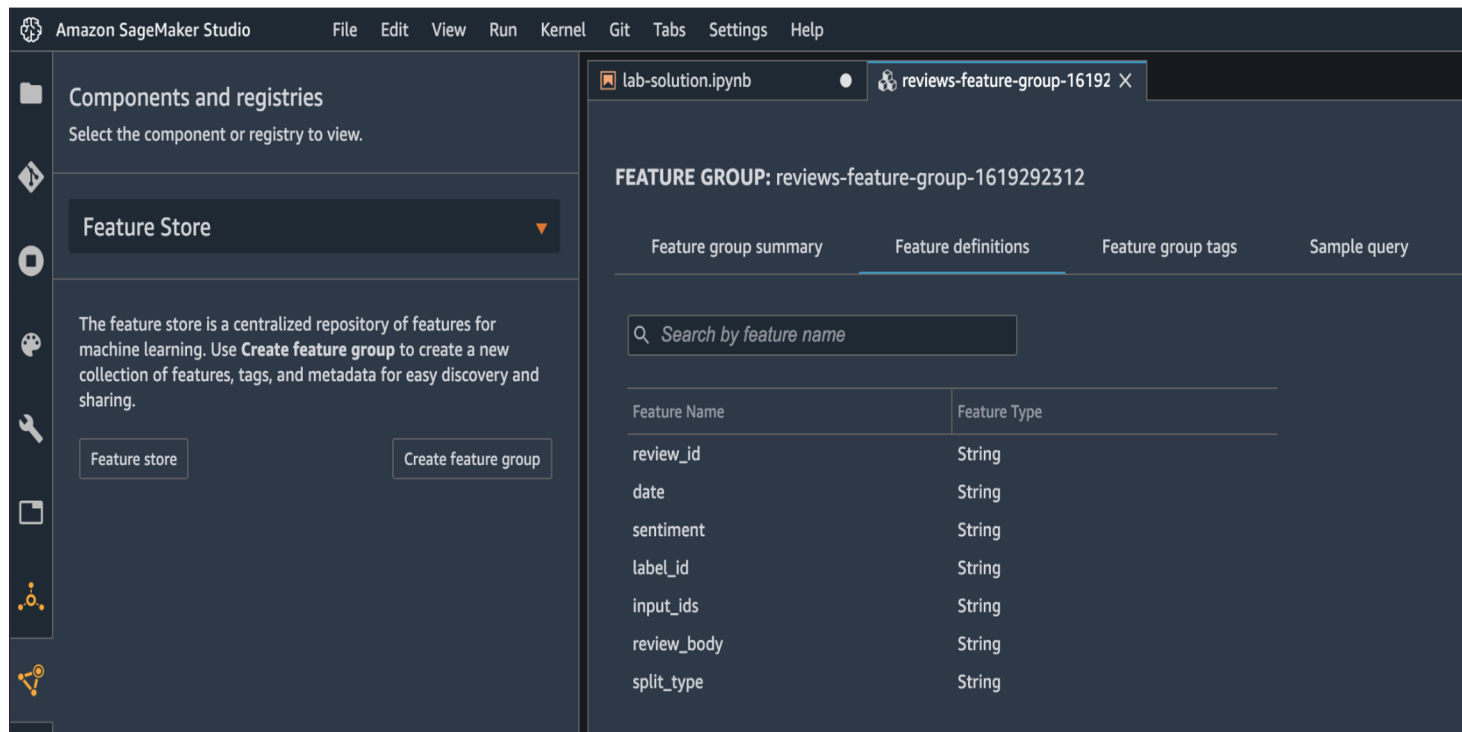
**Query string**

```
query_string = 'SELECT review_body, input_ids, input_mask, segment_ids,  
label_id FROM "{}" LIMIT 5'.format(reviews_feature_store_table)
```

```
reviews_feature_store_query.run(  
    query_string=..., ...)
```

**Execute  
the query**

# Amazon SageMaker Feature Store In SageMaker Studio



The screenshot displays the Amazon SageMaker Studio interface. On the left sidebar, under 'Components and registries', the 'Feature Store' is selected. The main panel shows the 'FEATURE GROUP: reviews-feature-group-1619292312' details. The 'Feature definitions' tab is active, displaying a table of features.

**Feature group summary** | **Feature definitions** | Feature group tags | Sample query

Search by feature name

Feature Name	Feature Type
review_id	String
date	String
sentiment	String
label_id	String
input_ids	String
review_body	String
split_type	String



# Amazon SageMaker Feature Store In SageMaker Studio

**FEATURE GROUP:** reviews-feature-group-1619921992

Feature group summary

Feature definitions

Feature group tags

Sample query

Use the buttons below to generate the query to perform the action with the feature data for this feature group. You can copy and paste this query to use with SageMaker Data Wrangler or in to any query interface for querying data from the offline store.

Interactive Exploration

Time travel

Remove tombstone

Remove duplicates

```
SELECT *  
FROM sagemaker_featurestore.reviews-feature-group-1619921992-1619922582  
LIMIT 1000
```

# Amazon SageMaker Feature Store In SageMaker Studio

	date	review_id	sentiment	label_id	input_ids	review_body
0	2021-04-29T18:34:07Z	14136	1	2	[0, 713, 16, 10, 182, 22, 4903, 3760, 254, 22, 2125, 4, 939, 657, 24, 328, 939, 2813, 6215, 74, ...	This is a very "retailer " piece. i love it! i wish retailer would bring back more pieces like t...
1	2021-04-29T18:34:07Z	4026	0	1	[0, 100, 1432, 5, 6173, 8, 2162, 10, 2514, 1836, 11, 5, 2440, 33953, 4, 939, 524, 2333, 10, 1836...	I followed the reviews and bought a larger size in the blue stripe. i am usually a size 8 but or...
2	2021-04-29T18:34:07Z	7522	-1	0	[0, 713, 8443, 16, 98, 11962, 8, 10698, 1969, 8, 939, 657, 5, 32847, 4, 959, 1437, 24, 24232, 90...	This jacket is so cute and fits perfect and i love the motif. however it deposited black linty ...
3	2021-04-29T18:34:07Z	7618	-1	0	[0, 133, 19111, 738, 8, 1421, 738, 32, 2198, 430, 4, 24, 18, 101, 45, 190, 5, 276, 3588, 4, 5, 5...	The catalog shot and model shot are completely different. it's like not even the same dress. the...
4	2021-04-29T18:34:07Z	11942	1	2	[0, 100, 269, 101, 5, 356, 9, 42, 8443, 1437, 53, 939, 206, 24, 1237, 650, 8, 939, 531, 671, 5, ...	I really like the look of this jacket but i think it runs small and i must return the one i rec...