

Copyright Notice

These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

For the rest of the details of the license, see <https://creativecommons.org/licenses/by-sa/2.0/legalcode>



DeepLearning.AI



Welcome to Course 3

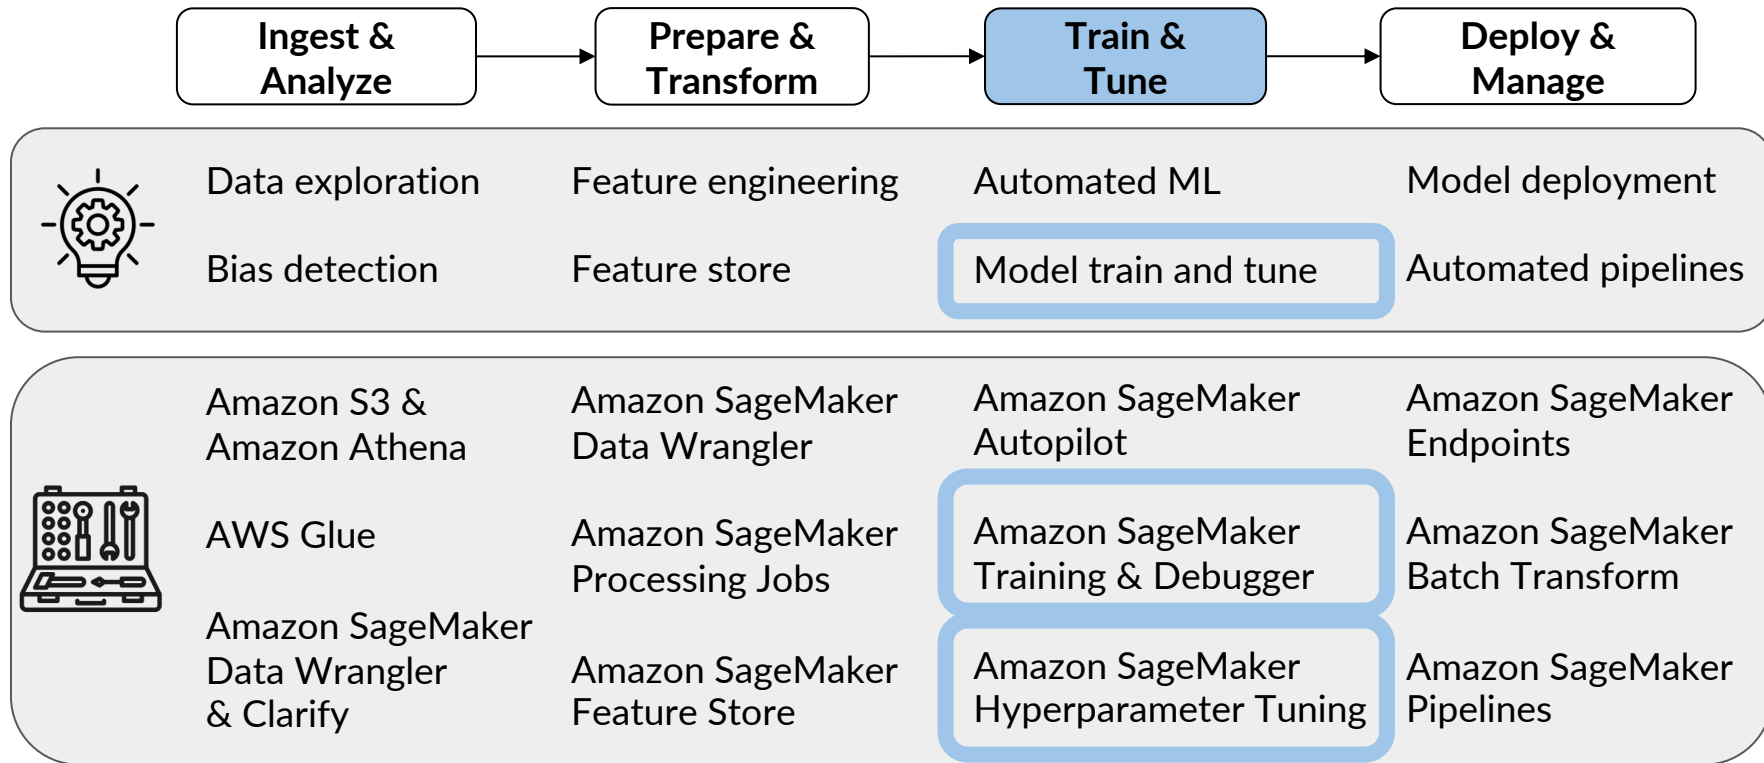


DeepLearning.AI



Advanced Model Training

Machine Learning Workflow



Model Tuning

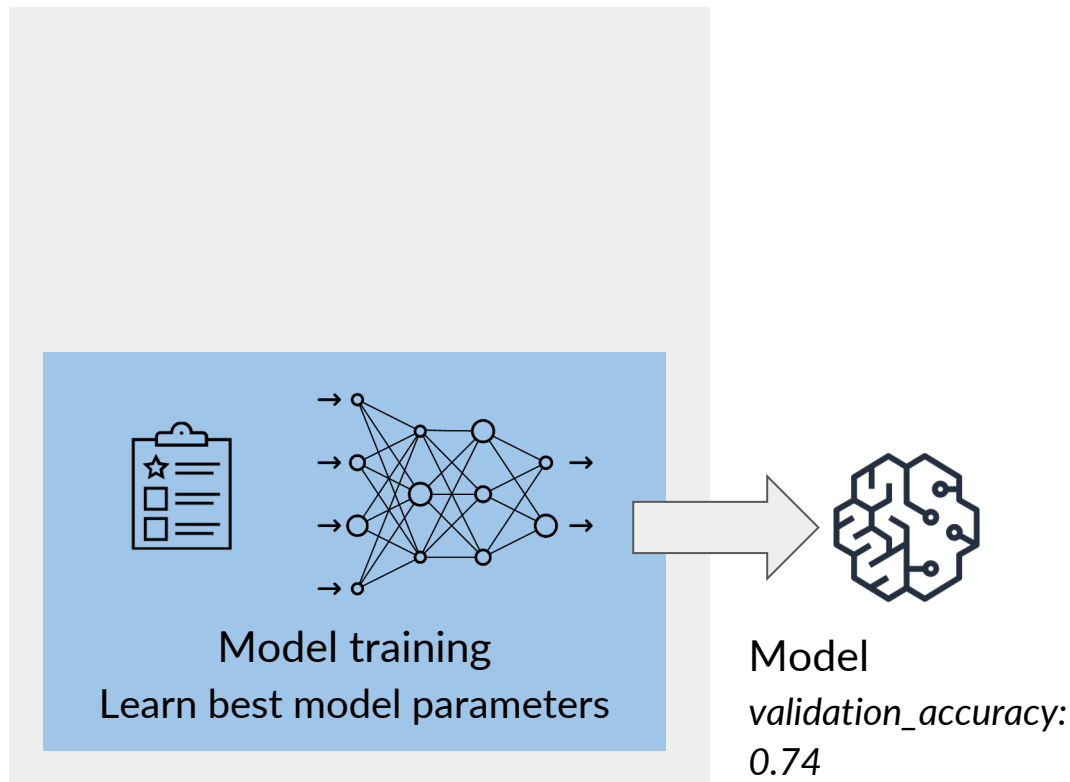


Model Tuning

Model parameters



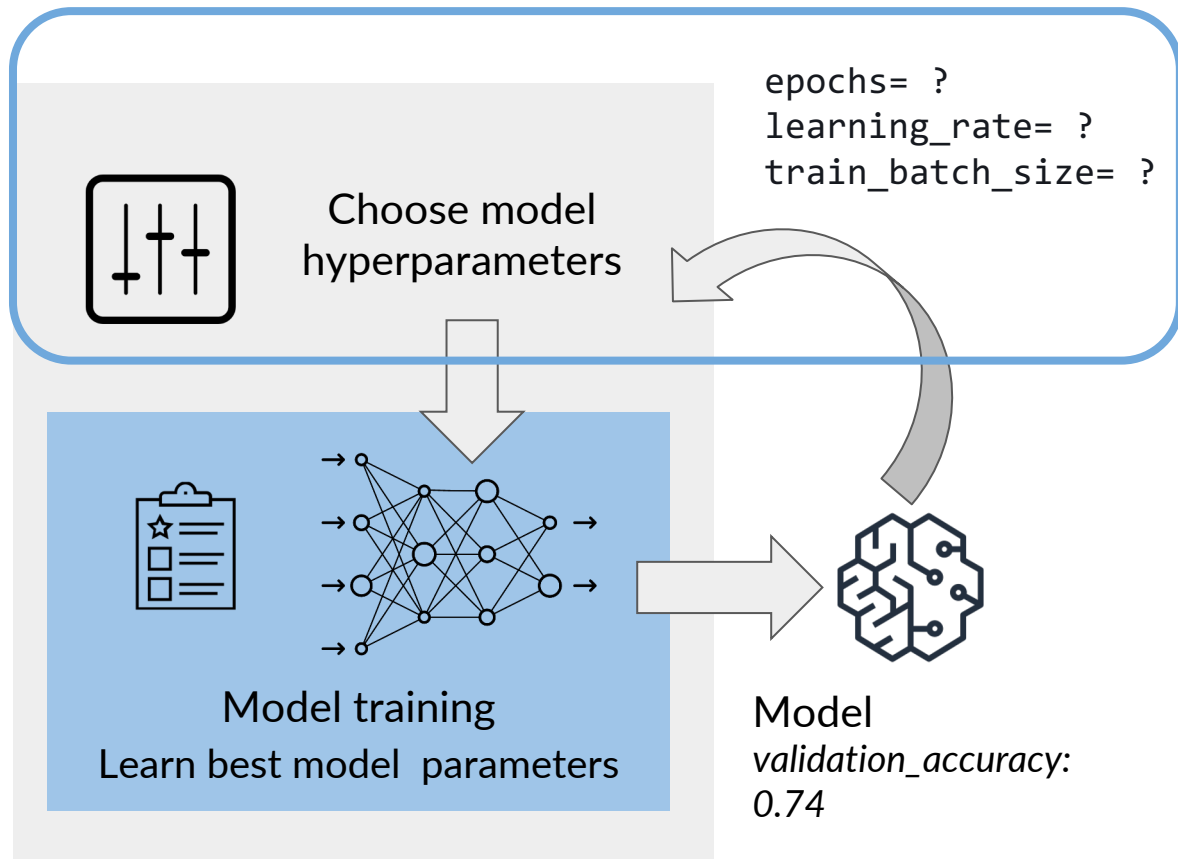
roberta-base
125M parameters



Model Tuning

Model hyperparameters

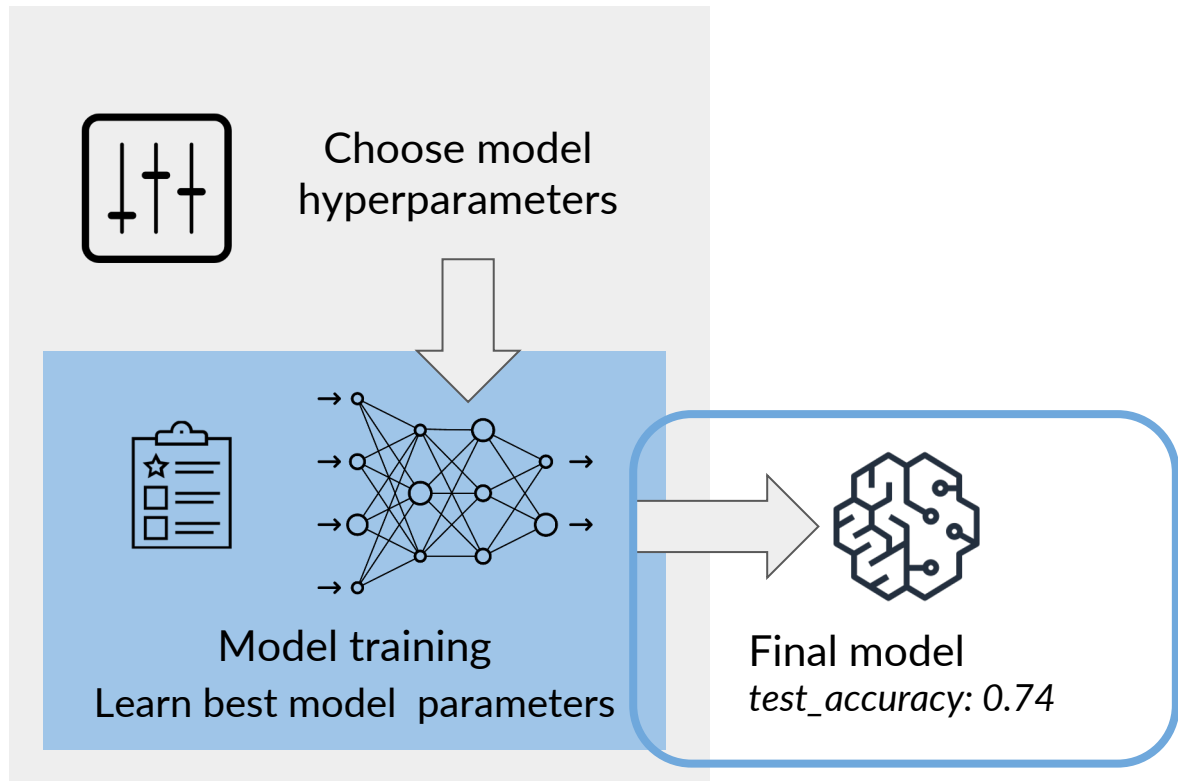
Model parameters



Model Evaluation

"If you can't measure it,
you can't improve it."


-- Peter Drucker



Manual vs. Automatic Model Tuning



Manual tuning



Automatic
model tuning

Popular Algorithms for Automatic Model Tuning





Automatic
model tuning

- Grid search
- Random search
- Bayesian optimization
- Hyperband



Grid Search

- Define sets of hyperparameters
- Test **every** combination
- Select the best performing hyperparameters



Grid Search



- Define sets of hyperparameters
- Test **every** combination
- Select the best performing hyperparameters



- + Explores all combinations
- + Works for small number of parameters
- Time-consuming
- Doesn't scale to large numbers of parameters



Random Search

- Define sets of hyperparameters
 - Define search space & stop criteria
 - Test **random** combinations within search space
 - Select the best performing hyperparameters
-



Random Search



- Define sets of hyperparameters
 - Define search space & stop criteria
 - Test **random** combinations within search space
 - Select the best performing hyperparameters
-



- + Faster compared to grid search
- Might miss better performing hyperparameters



Bayesian Optimization

- Treat HPT like a **regression** problem (surrogate model)
- Start from random hyperparameters
- Narrow down search space around better performing hyperparameters




Bayesian Optimization

- Treat HPT like a **regression** problem (surrogate model)
- Start from random hyperparameters
- Narrow down search space around better performing hyperparameters




- + More efficient in finding best hyperparameters
- Requires sequential execution
- Might get stuck in local minima




When you use gradient descent to minimize the Loss Function, it is a possibility that this algorithm might get stuck with a Local minima of the Loss Function and might not find the Global minima.

Hyperband



- **Bandit-based** approach
- Start from random hyperparameters
- Explore sets of hyperparameters for few iterations
- Choose best and explore longer
- Repeat until `max_iterations` reached or one candidate left




Bandit approaches typically use a combination of exploitation and exploration to find the best possible hyperparameters. The strength of the bandit approaches is that dynamic pull between exploitation and exploration.


You start with the larger space of random hyperparameter set and then you explore a random subset of these hyperparameters for a few iterations. After the first few iterations, you discard the worst performing half of the hyperparameter sets. In the subsequent few iterations, you continue to explore the best performing hyperparameters from the previous iteration. You continue this process until the set time is elapsed or you remain with just one possible candidate.

Hyperband clearly stands out by spending the time much more efficiently than other approaches we discussed to explore the hyperparameter values using the combination of exploitation and exploration. On the downside, it might discard good candidates very early on and these could be the candidate that converge slowly.

Hyperband



- **Bandit-based** approach
- Start from random hyperparameters
- Explore sets of hyperparameters for few iterations
- Choose best and explore longer
- Repeat until `max_iterations` reached or one candidate left

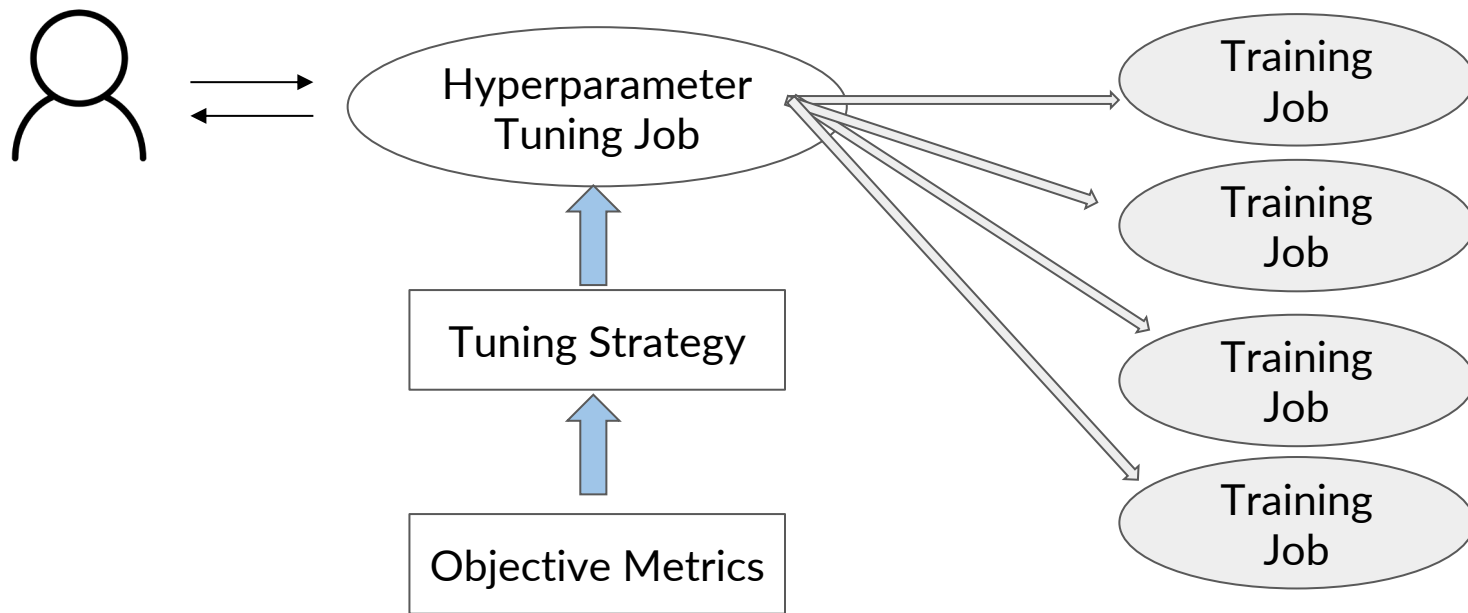


- + Spends time efficiently (explore-exploit theory)
- Might discard good candidates early that converge slowly

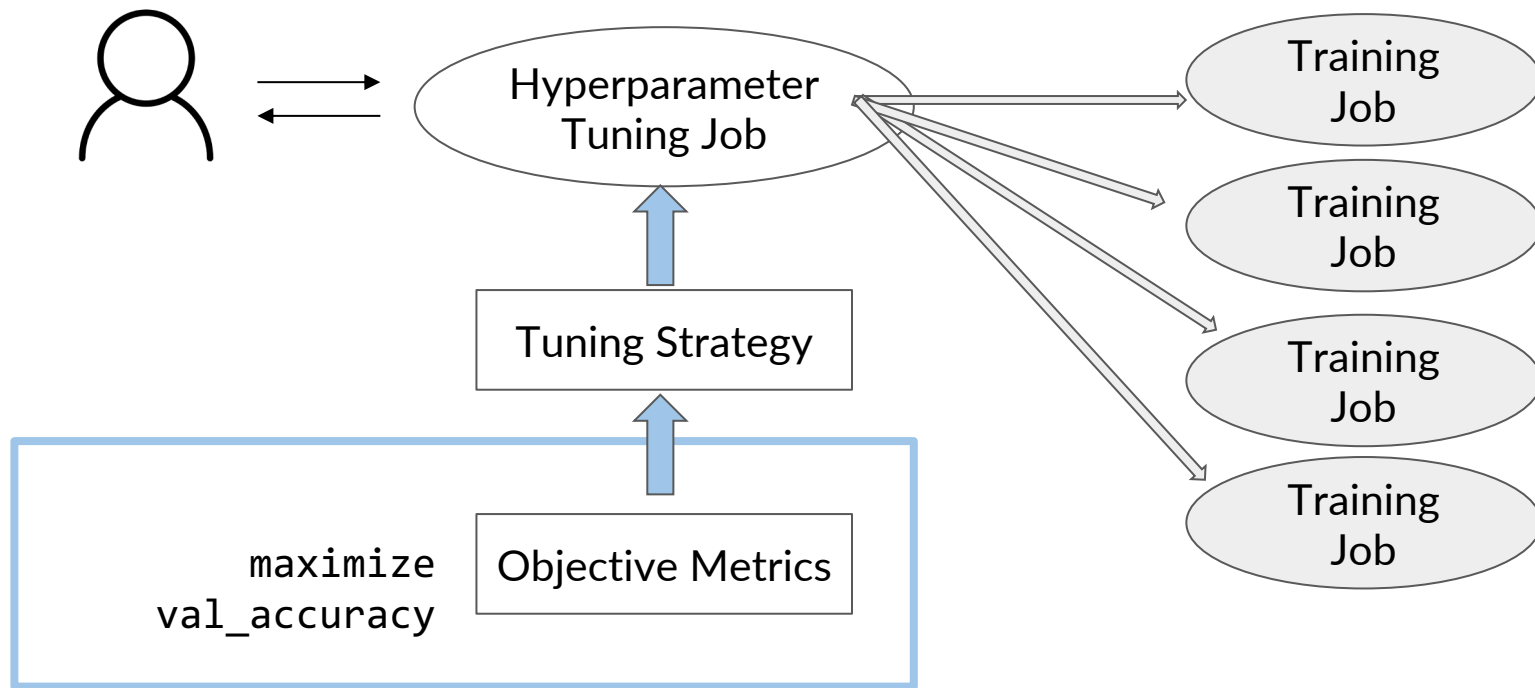
Tune a BERT- based Text Classifier



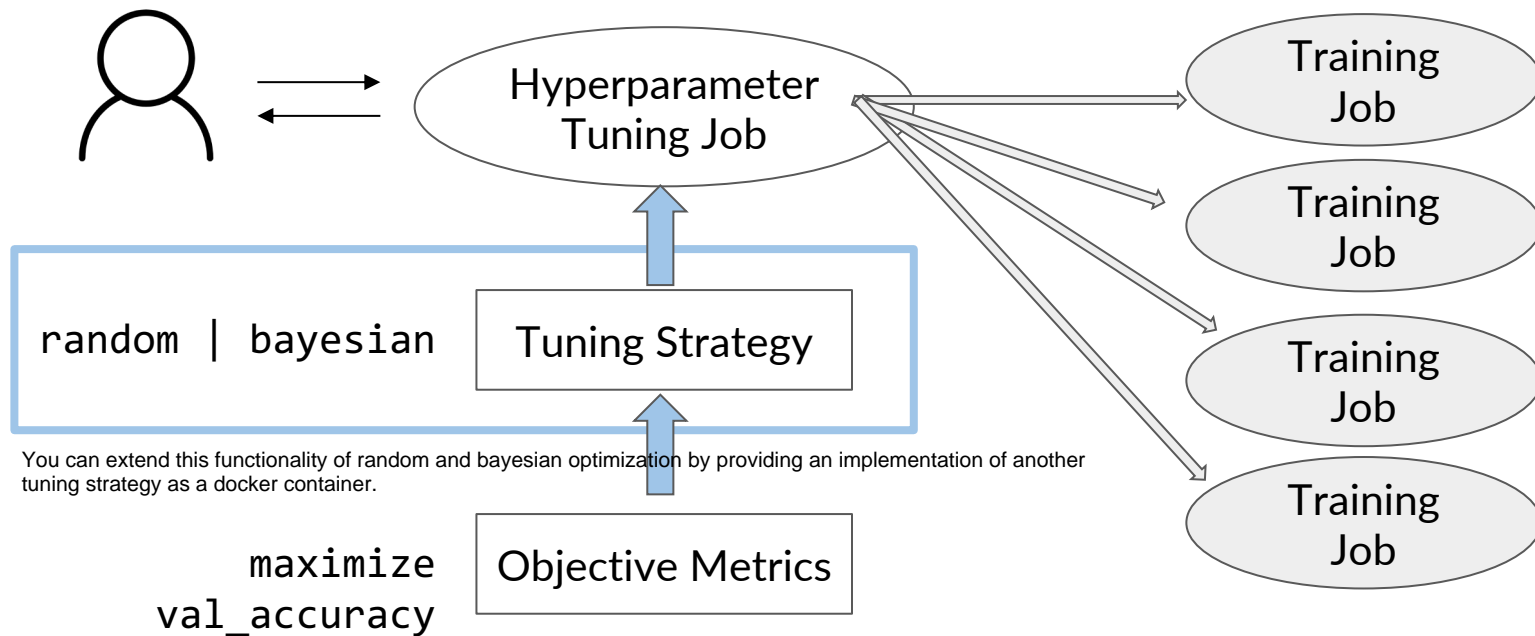
Amazon SageMaker Hyperparameter Tuning (HPT)



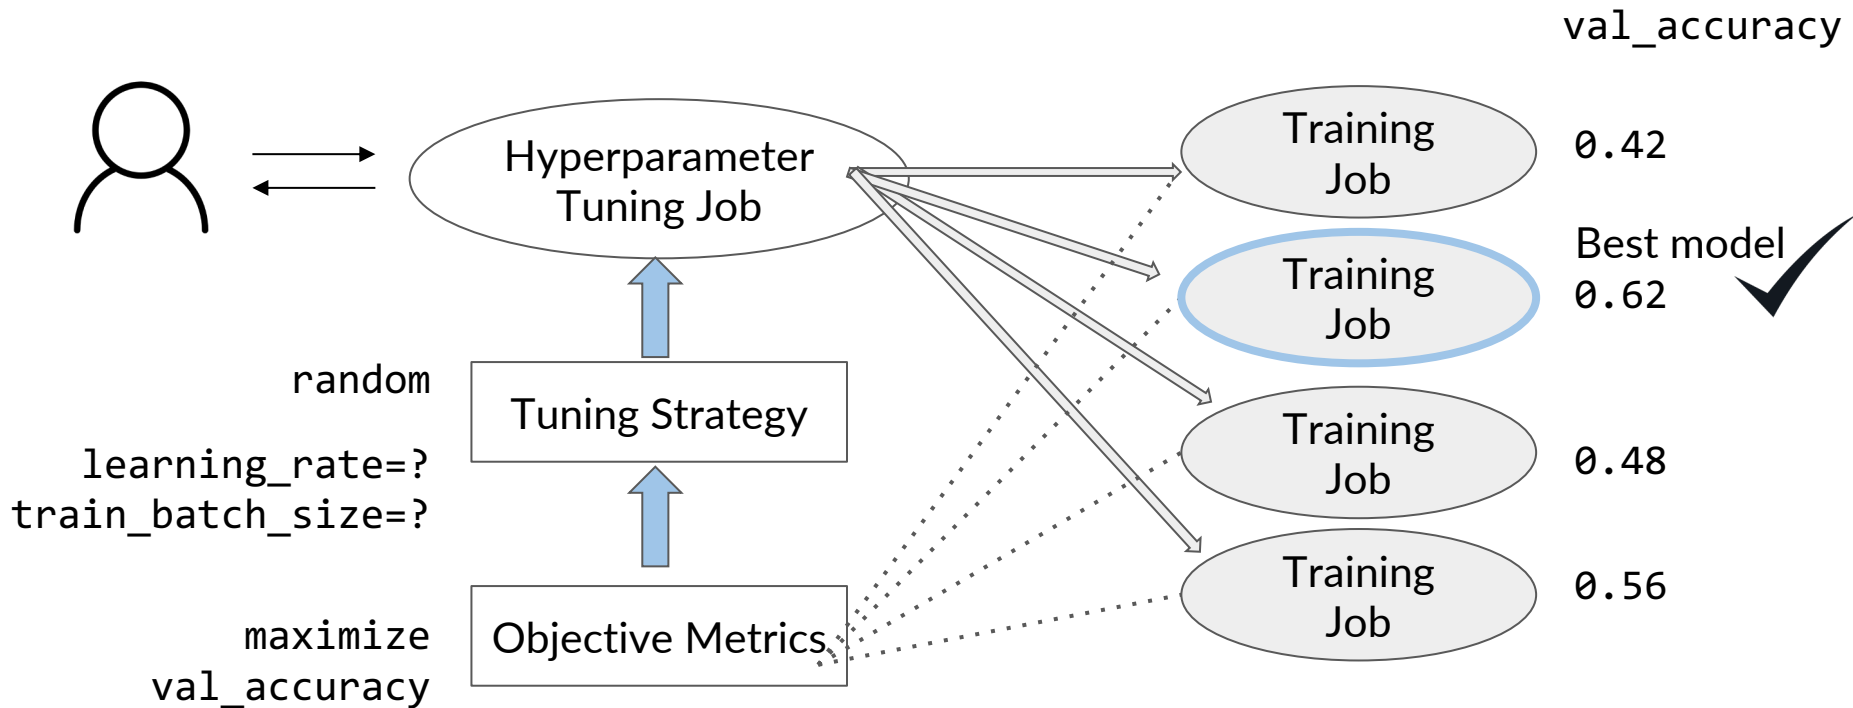
Amazon SageMaker Hyperparameter Tuning (HPT)



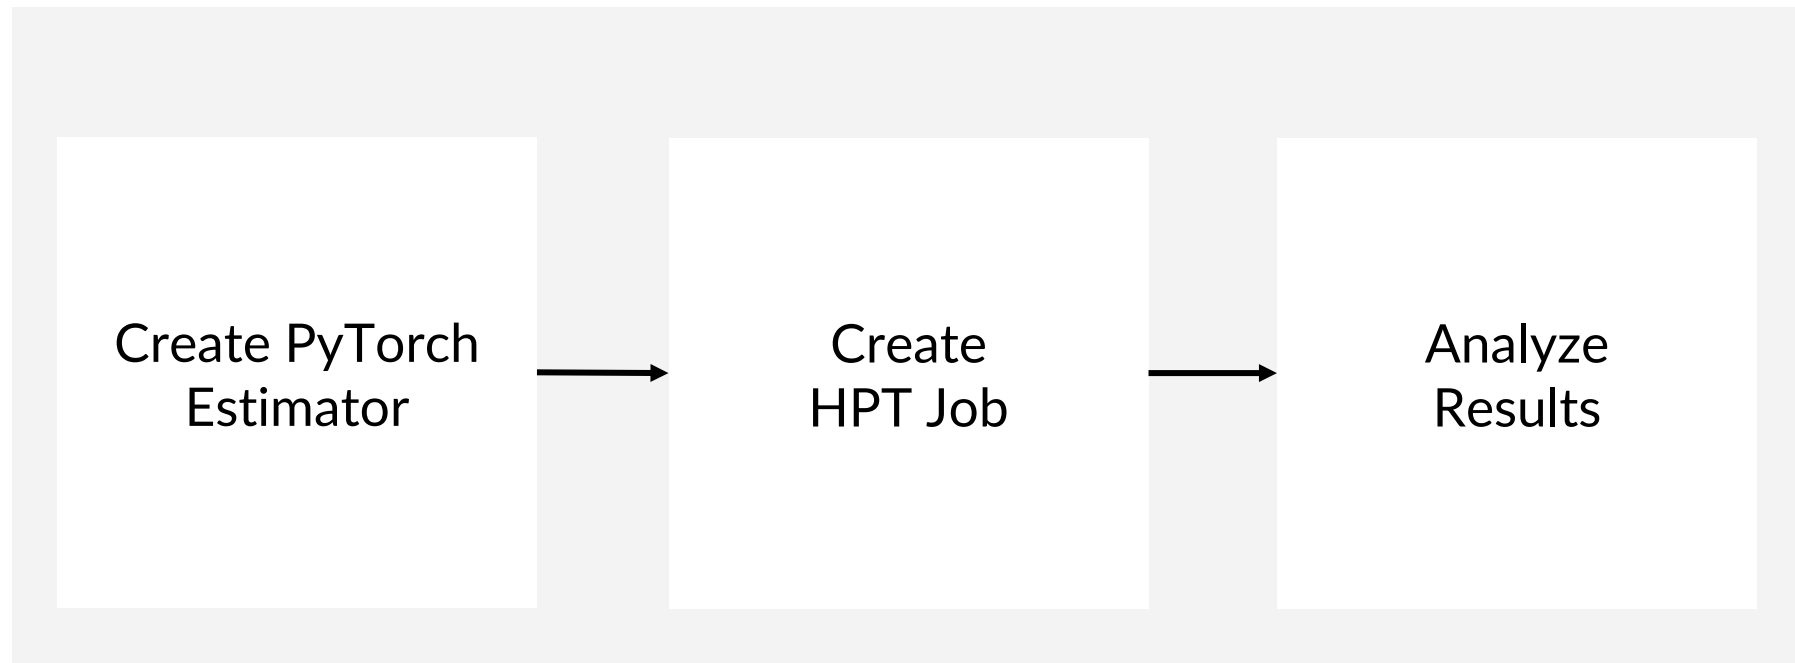
Amazon SageMaker Hyperparameter Tuning (HPT)



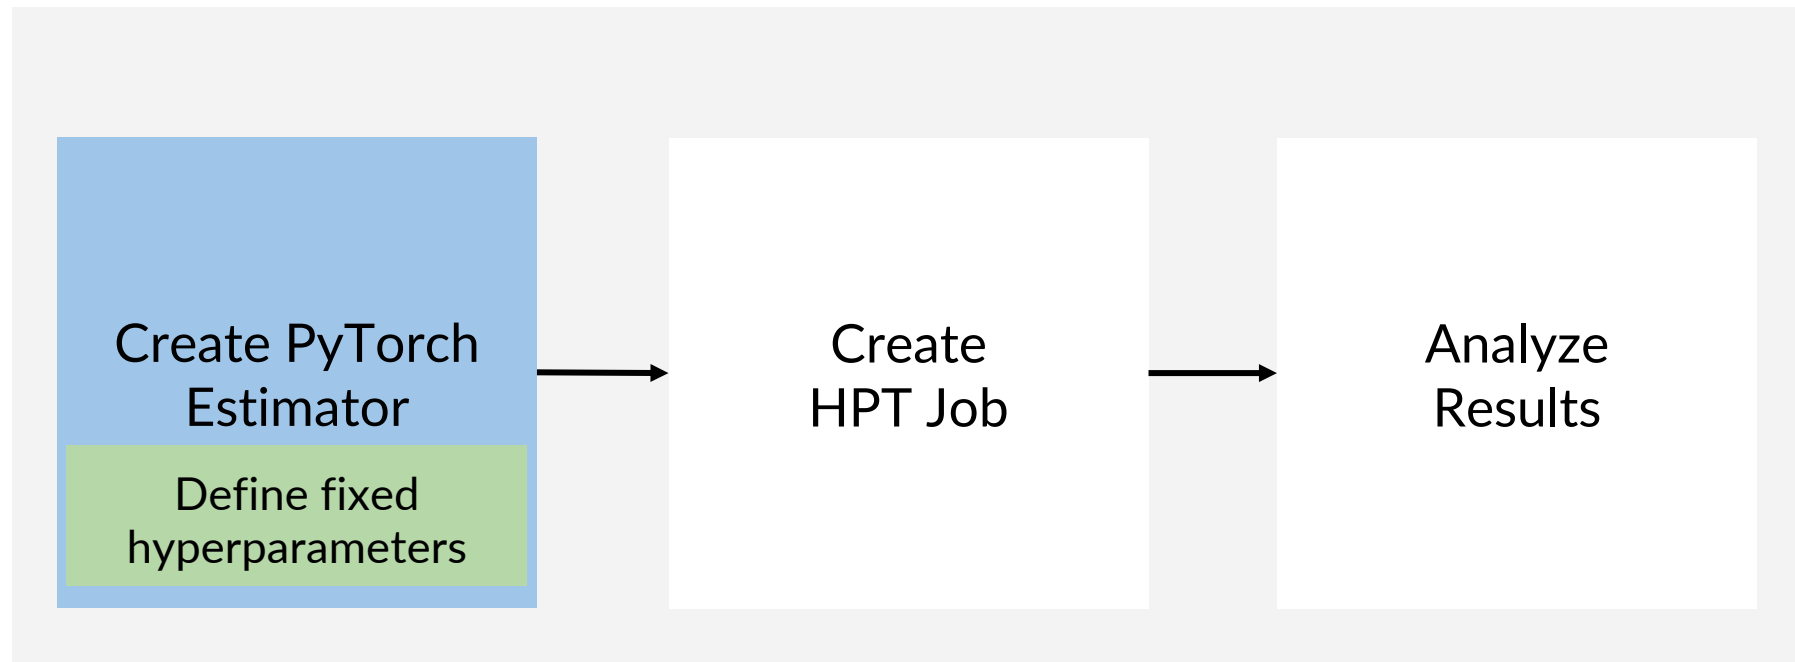
Tune BERT Text Classifier



Steps



Steps



Define Fixed Hyperparameters

```
hyperparameters={  
    'epochs': 3,  
    'train_steps_per_epoch': 50,  
    'validation_batch_size': 64,  
    'validation_steps_per_epoch': 50,  
    'freeze_bert_layer': False,  
    'seed': 42,  
    'max_seq_length': 64,  
    'backend': 'gloo',  
    'run_validation': True,  
    'run_sample_predictions': False  
}
```

Create PyTorch estimator

Define fixed
hyperparameters

Create PyTorch Estimator

Create PyTorch estimator

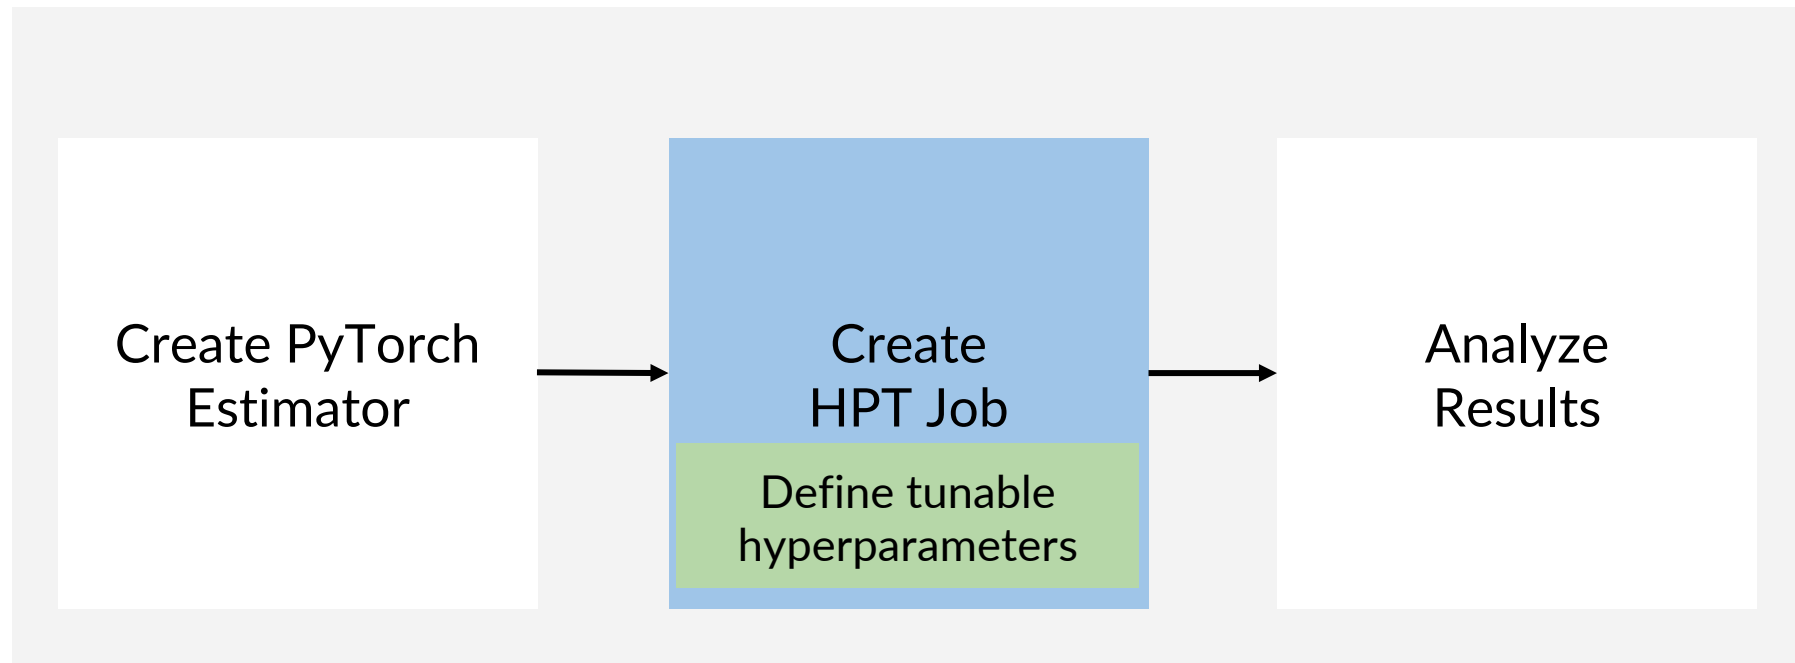
Define fixed
hyperparameters

```
from sagemaker.pytorch import PyTorch as PyTorchEstimator
```

```
estimator = PyTorchEstimator(  
    entry_point='train.py',  
    ...  
    hyperparameters=hyperparameters,  
  
)
```

Fixed hyperparameters
in estimator

Steps



Define Tunable Hyperparameters

Create HPT job

Define tunable
hyperparameters

```
from sagemaker.tuner import CategoricalParameter
from sagemaker.tuner import ContinuousParameter
from sagemaker.tuner import IntegerParameter
```

Specify parameters

```
hyperparameter_ranges = {
    'learning_rate': ContinuousParameter(0.00001, 0.00005,
scaling_type='Linear'),
    'train_batch_size': CategoricalParameter([128, 256]),
}
```

Specify hyperparameter types

Specify ranges

How to Choose Hyperparameter Types

Categorical

'train_batch_size':

CategoricalParameter([128, 256])

'freeze_bert_layer':

CategoricalParameter([True, False])

How to Choose Hyperparameter Types

Categorical

'train_batch_size':

CategoricalParameter([128, 256])

'freeze_bert_layer':

CategoricalParameter([True, False])

Integer

'train_batch_size':

IntegerParameter(16, 1024, scaling_type='Logarithmic')

'Logarithmic scales works well if you want to explore large ranges quickly.'

If you need to
explore large
ranges quickly

How to Choose Hyperparameter Types

Categorical

'train_batch_size':

CategoricalParameter([128, 256])

'freeze_bert_layer':

CategoricalParameter([True, False])

Integer

'train_batch_size':

IntegerParameter(16, 1024, scaling_type='Logarithmic')

Continuous

'learning_rate':

ContinuousParameter(0.00001, 0.00005, scaling_type='Linear')

Create Amazon SageMaker HPT job

Create HPT job

Define tunable
hyperparameters

```
from sagemaker.tuner import HyperparameterTuner
```

```
tuner = HyperparameterTuner(  
    estimator=...,  
    hyperparameter_ranges=...,  
    objective_type=...,  
    objective_metric_name=...,  
    strategy=...,  
)
```

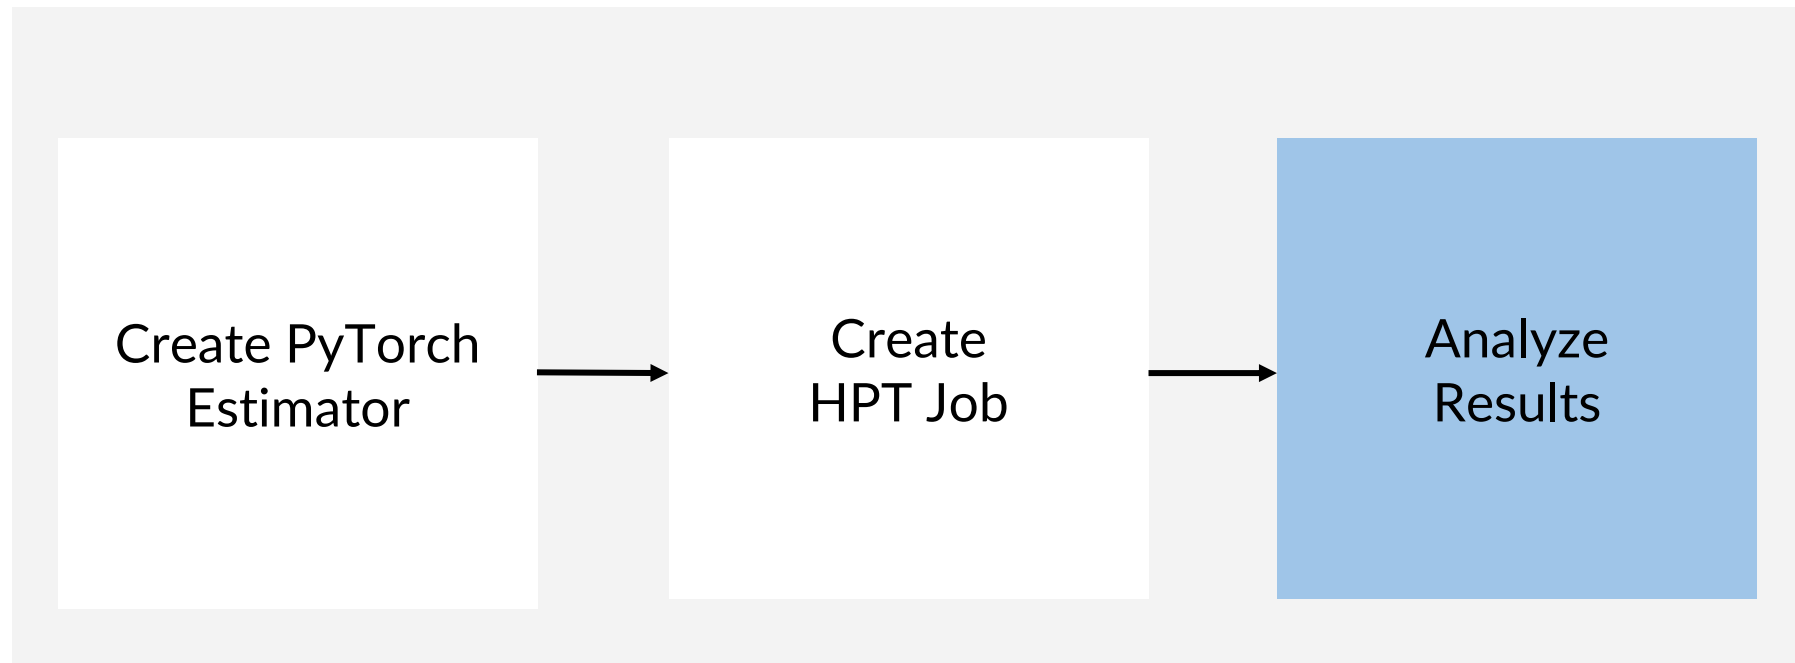
Pass in estimator

Configure
Hyperparameter ranges

Run HPT job with .fit()

```
tuner.fit(inputs={...}, ...)
```

Steps



Analyze Results

Analyze Results


```
df_results = tuner.analytics().dataframe()
```

Analyze Results

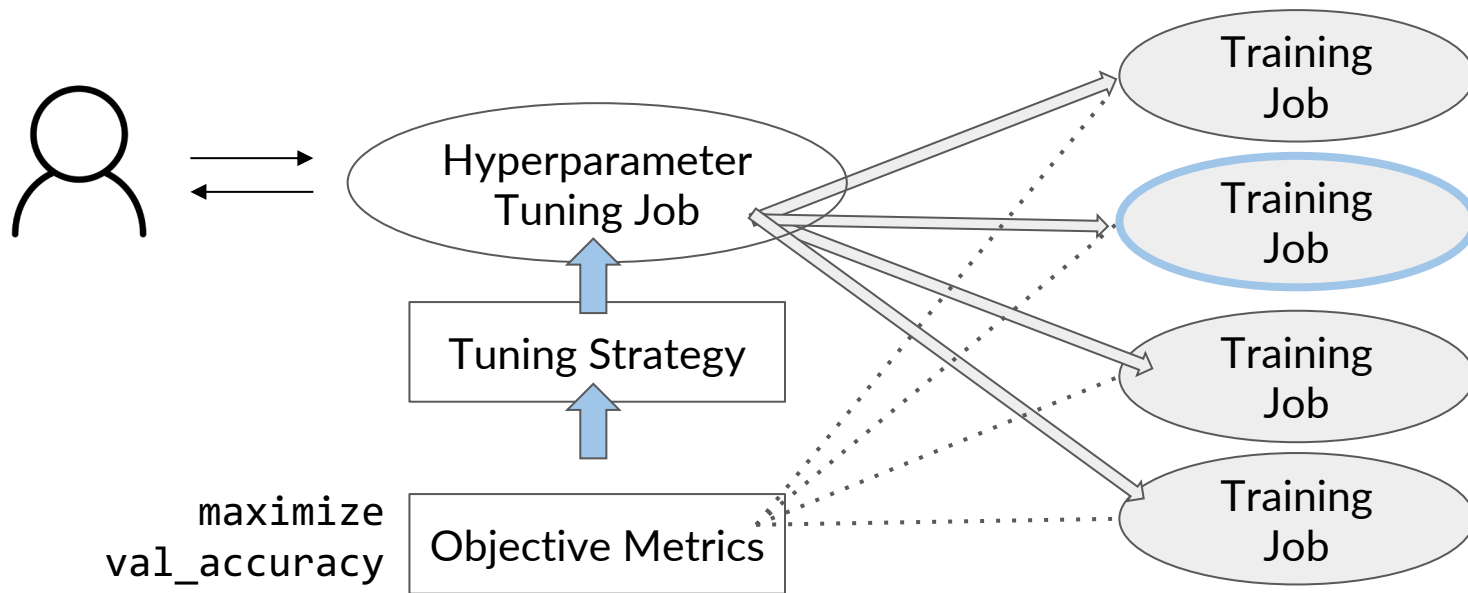
Analyze Results

```
df_results = tuner.analytics().dataframe()
```

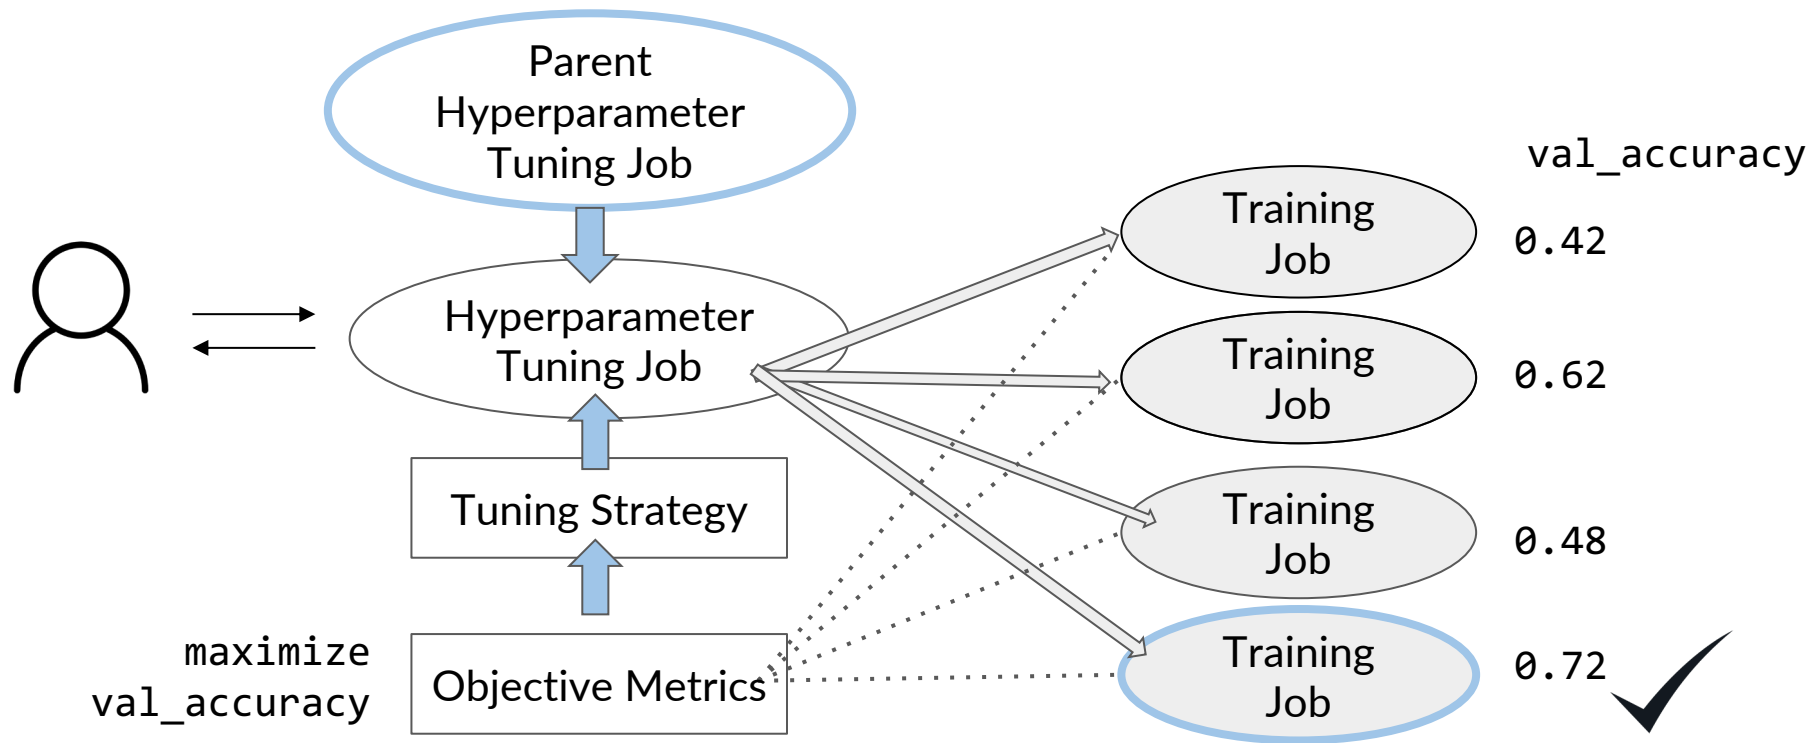
learning_rate	train_batch_size	TrainingJobName	TrainingJobStatus	FinalObjectiveValue
0.000021	"128"	pytorch-training-210225-1535-001-71394bc3	Completed	44.939999
0.000035	"128"	pytorch-training-210225-1535-002-cf437bad	Completed	41.580002



Warm Start HPT Job



Warm Start HPT Job



Warm Start HPT Job

- IDENTICAL_DATA_AND_ALGORITHM
 - Same input data and training data
 - Update hyperparameter tuning ranges and maximum number of training jobs
- TRANSFER_LEARNING
 - Updated training data and different version of training algorithm

Configure Warm Start

```
from sagemaker.tuner import WarmStartConfig
from sagemaker.tuner import WarmStartTypes
```

```
warm_start_config = WarmStartConfig(
    warm_start_type=WarmStartTypes.IDENTICAL_DATA_AND_ALGORITHM,
    parents=<PARENT_TUNING_JOB_NAME>)
```

IDENTICAL_DATA_AND_ALGORITHM
or TRANSFER_LEARNING

```
tuner = HyperparameterTuner(
    ...
    warm_start_config=warm_start_config)
```

Specify parent tuning job


```
tuner.fit(...)
```

Pass warm start config
in HyperparameterTuner

Best Practices - SageMaker HyperParameter Tuning

Hyperparameter Tuning is a time and computation intensive task.

The computational complexity is directly proportional to the number of hyperparameters that you tune.



- Select a small number of hyperparameters
- Select a small range for hyperparameters
- Enable warm start
- Enable early stop to save tuning time and costs
- Select a small number of concurrent training jobs

On one hand, if you use a larger number of concurrent jobs, the tuning process will be completed faster. But in fact, the hyperparameter tuning process is able to find best possible results only by depending on the previously completed training jobs.

Best Practices - Monitoring Training Resources

- Right size compute resources
- Requires empirical testing
- Amazon CloudWatch Metrics
- Insights from Amazon SageMaker Debugger

Checkpointing



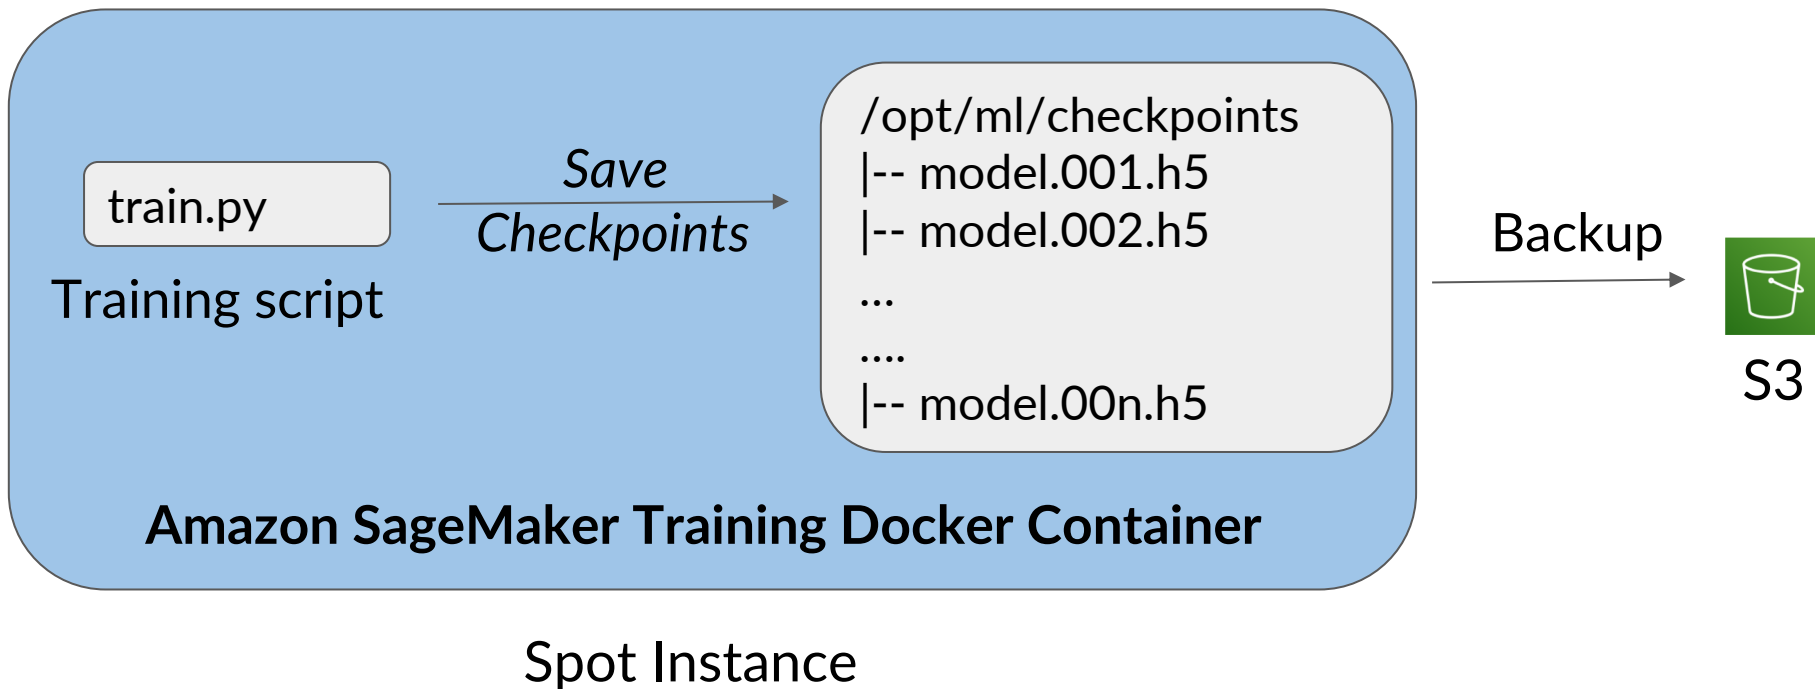
Machine Learning Checkpointing

- Save state of ML models during training
- Checkpoints : Snapshots of the model
 - Model architecture which allows you to recreate the model training once it stopped
 - Model weights that have been learned in the training process so far
 - Training configurations training configuration such as number of epochs that have been executed, and the optimizer used, and the loss observed so far in training, and other metadata information
 - Optimizer This optimizer state allows you to easily resume the training job from where it has stopped.
- Frequency and number of checkpoints

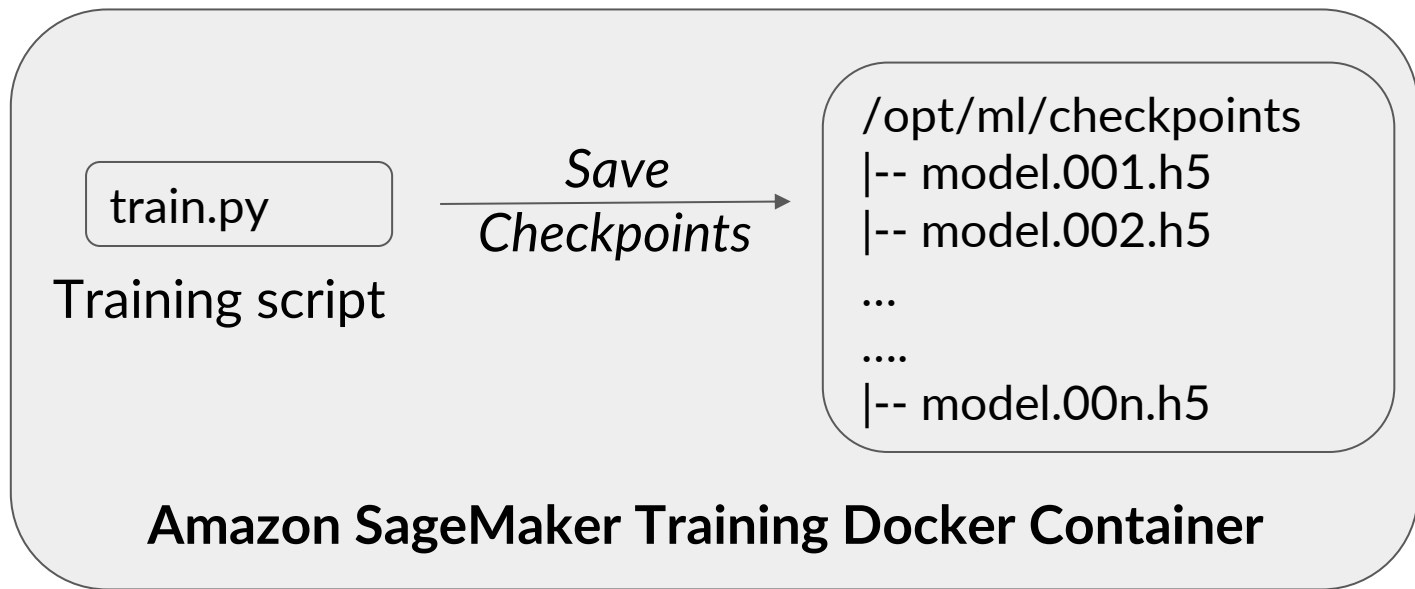
When configuring your new training job with checkpointing take two things into consideration, one is the frequency of checkpointing, and the second is the number of checkpoint files you are saving each time.

Amazon SageMaker Managed Spot

Managed Spot capability allows you to save training costs. Managed Spot is based on the concept of Spot Instances that offer speed and unused capacity to users at discount prices.



Amazon SageMaker Managed Spot



S3

Spot Instance (Terminated)

Amazon SageMaker Spot Training

Spot Instance (Terminated)

train.py

Training script

*Save
Checkpoints*

/opt/ml/checkpoints

|-- model.001.h5

|-- model.002.h5

...

....

|-- model.00n.h5

Amazon SageMaker Training Docker Container

Spot Instance (New)



S3

Sync

Distributed Training Strategies



Challenges



Increased training data volume




Increased model size and complexity


Distributed Training

The training load is split across multiple CPUs and GPUs, also called as devices within a single Compute Node.

Or the node can be distributed across multiple compute nodes or compute instances that form a compute cluster.




Compute Node



Compute Cluster

Distributed Training Strategies





Data Parallelism



Model Parallelism

Distributed Training Strategies - Data Parallelism

- Training data split up
- Model replicated on all nodes



Distributed Training Strategies - Model Parallelism

- Training data replicated
- Model split up on all nodes



Amazon SageMaker Estimator

```
from sagemaker.pytorch import PyTorch
estimator = PyTorch(
    entry_point='train.py',
    role=sagemaker.get_execution_role(),
    framework_version='1.6.0',
    py_version='py3',
    instance_count=3,
    instance_type='ml.p3.16xlarge',
    distribution={
        'smdistributed': {
            'dataparallel': {
                'enabled': True
            }
        }
    }
)
estimator.fit()
```

Data Parallelism

Distribution Strategy

Amazon SageMaker Estimator


```
from sagemaker.pytorch import PyTorch
estimator = PyTorch(
    entry_point='train.py',
    role=sagemaker.get_execution_role(),
    framework_version='1.6.0',
    py_version='py3',
    instance_count=3,
    instance_type='ml.p3.16xlarge',
    distribution={'smdistributed':{'modelparallel':{'enabled': True}}}
)
estimator.fit()
```

Model Parallel
Distribution Strategy

Model Parallelism

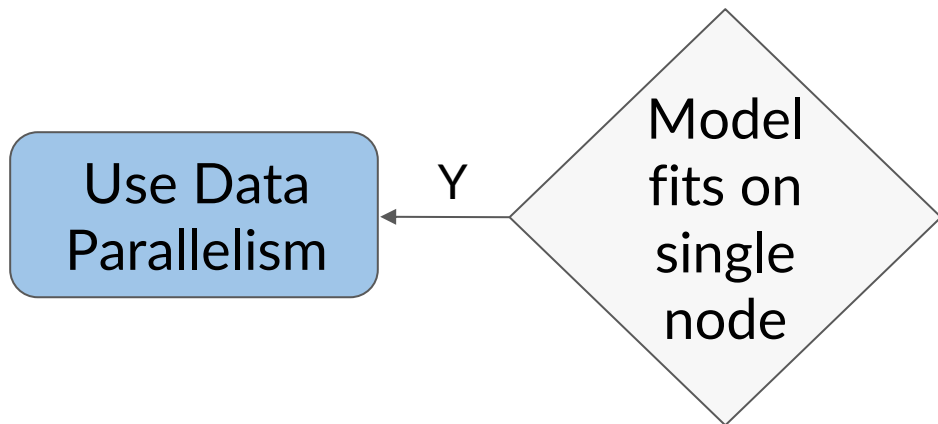
Choosing a Distribution Strategy

When choosing a distributed training strategy always keep in mind that if your training across multiple nodes or multiple instances, there is always a certain training overhead. The training overhead comes in the form of internode communication because of the data that needs to be exchanged between the multiple nodes of the cluster.




Choosing a Distribution Strategy


If the train model can fit on a single node's memory, then use data parallelism. In the situations where the model cannot fit on a single node's memory, you have some experimentation to do to see if you can reduce the model size to fit on that single node. All of these experimentations will include an effort to resize the model.




Choosing a Distribution Strategy




Choosing a Distribution Strategy




Choosing a Distribution Strategy




Choosing a Distribution Strategy



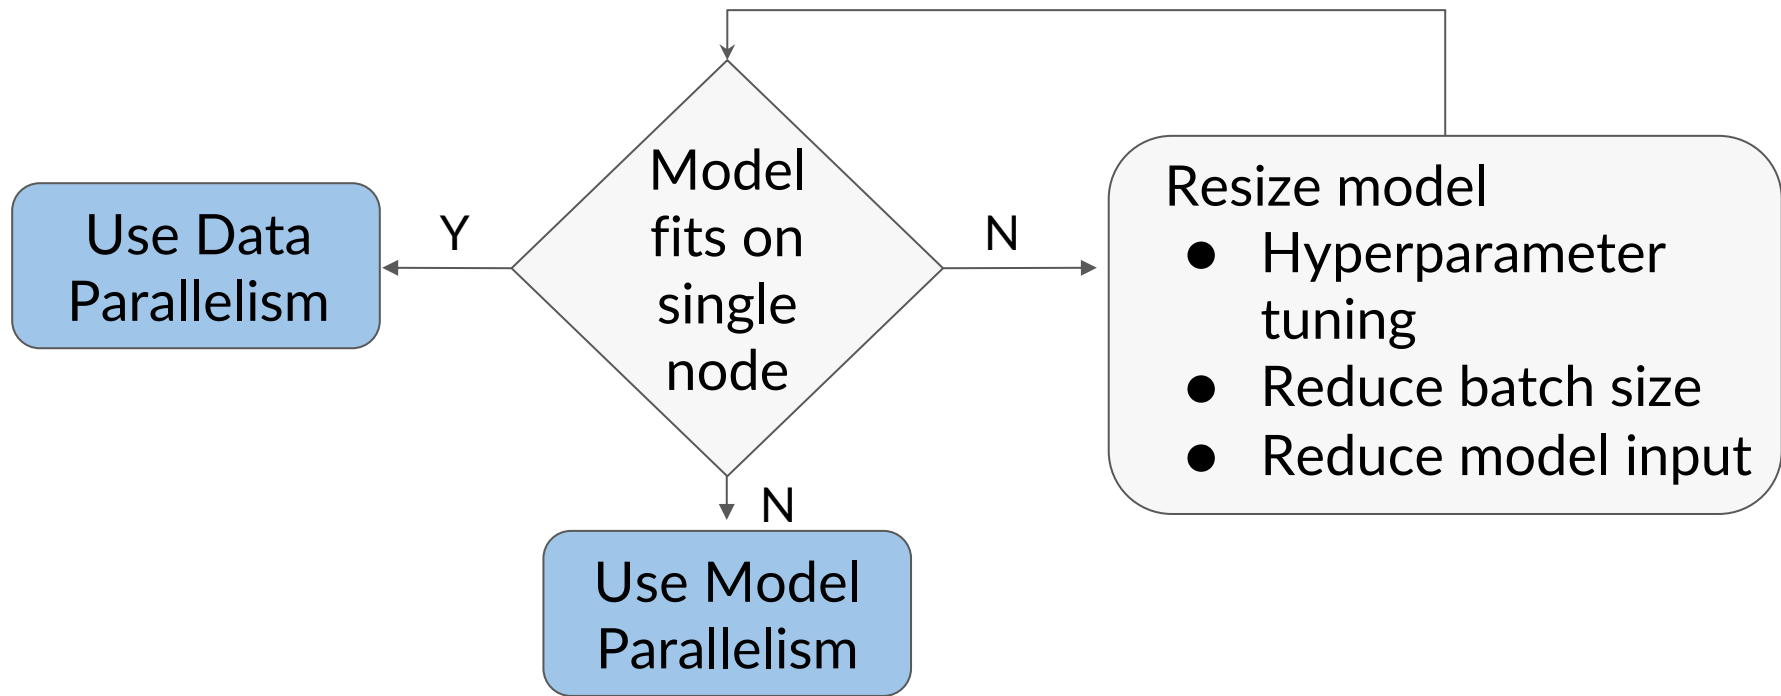
Choosing a Distribution Strategy



Choosing a Distribution Strategy



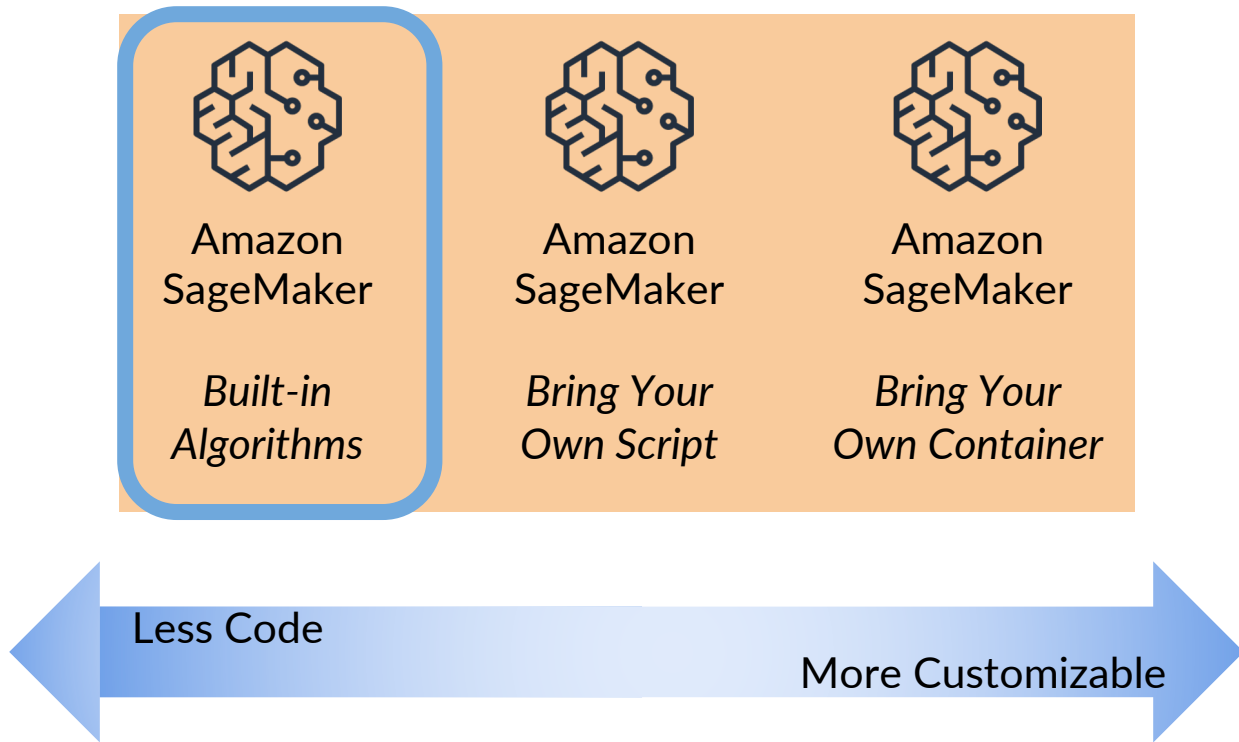
Choosing a Distribution Strategy



Custom Algorithms with Amazon SageMaker



Options on Amazon SageMaker

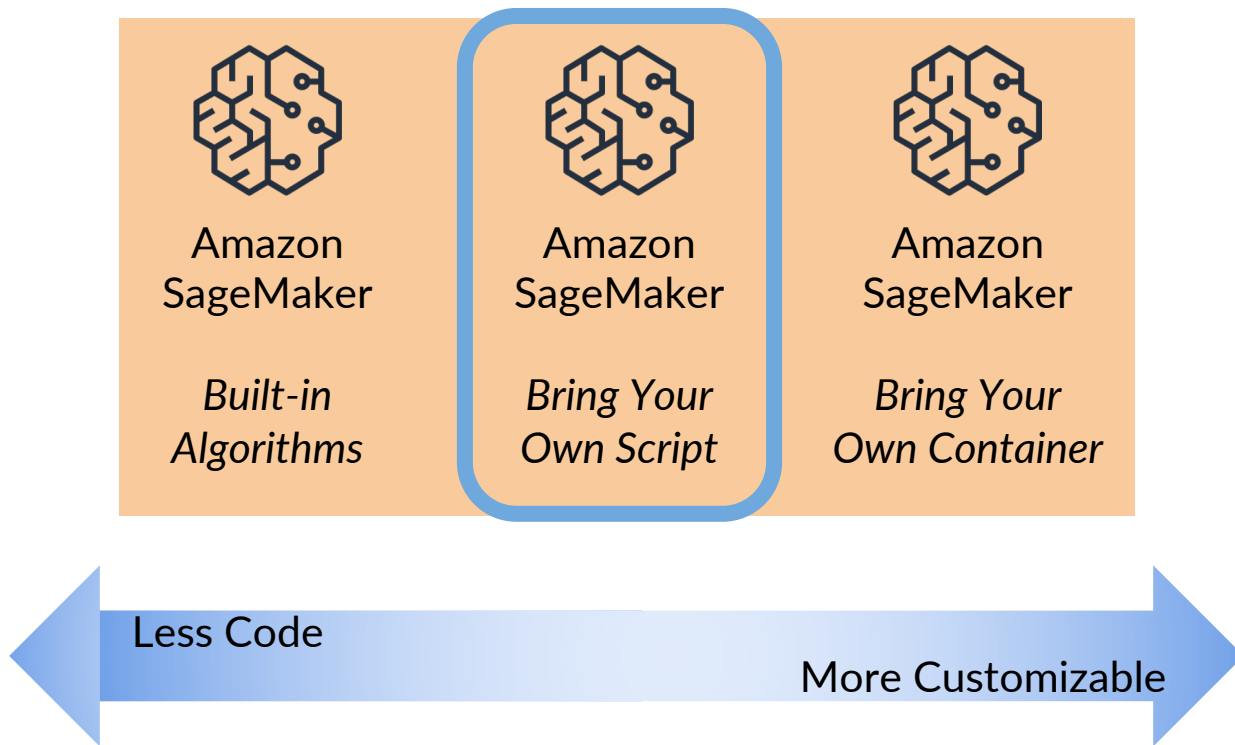


Amazon SageMaker Estimator

```
estimator =  
sagemaker.estimator.Estimator(image_uri=image_uri, ...)  
estimator.set_hyperparameters(...)  
estimator.fit(...)
```

Built-In Algorithms

Options on Amazon SageMaker



Amazon SageMaker Estimator

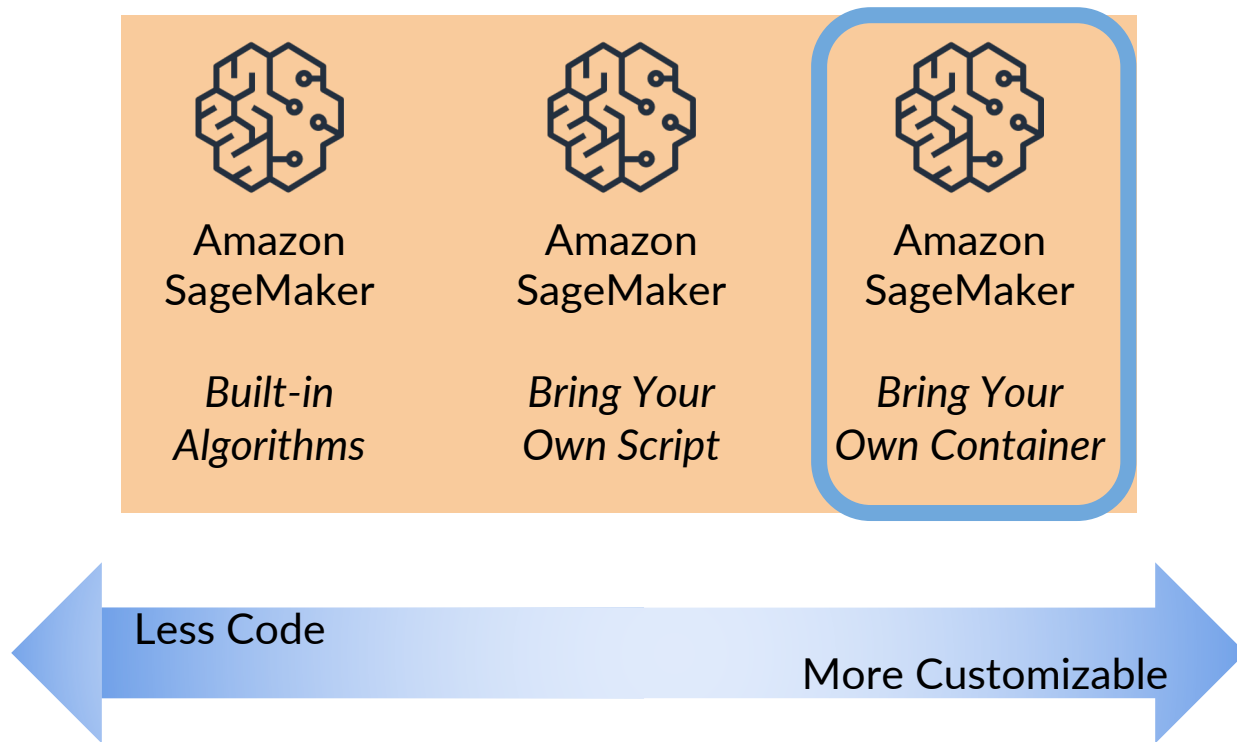
```
estimator =  
sagemaker.estimator.Estimator(image_uri=image_uri, ...)  
estimator.set_hyperparameters(...)  
estimator.fit(...)
```

Built-In Algorithms

```
from sagemaker.pytorch import PyTorch  
pytorch_estimator = PyTorch(  
    entry_point='train.py',  
    ...  
)
```

**Script Mode
PyTorch Container**

Training Options on Amazon SageMaker



Bring Your Own Container - Steps

Code



Containerize

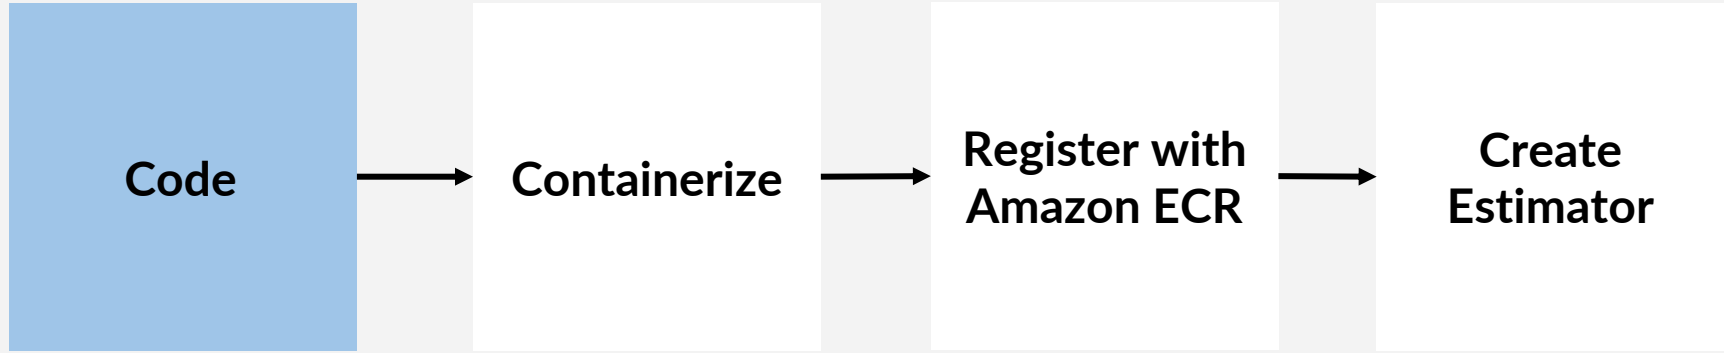


**Register with
Amazon ECR**

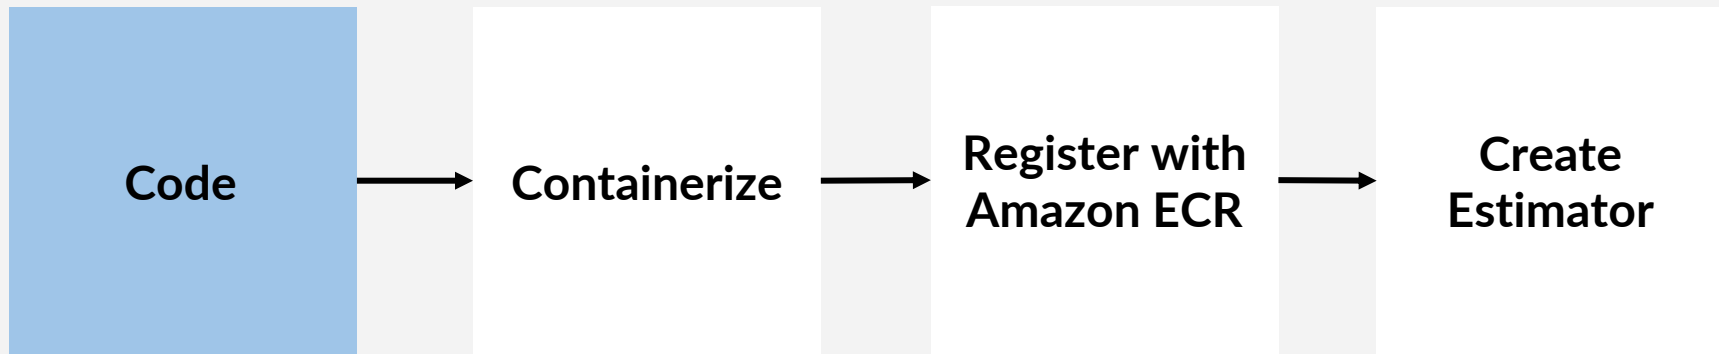


**Create
Estimator**

Bring Your Own Container - Steps



Bring Your Own Container - Steps



- Algorithm
- Training
- Inference

Bring Your Own Container - Steps

Code



Containerize

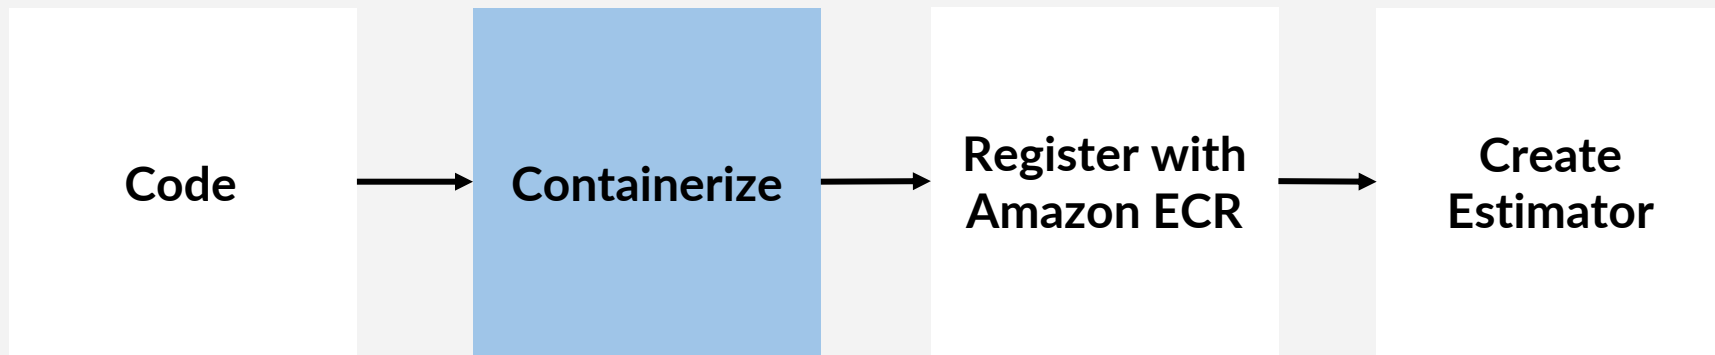


Register with
Amazon ECR



Create
Estimator

Bring Your Own Container - Steps



- `algorithm_name=tf-custom-container-test`
- `docker build -t ${algorithm_name} .`

Bring Your Own Container - Steps

Code



Containerize



Register with
Amazon ECR



Create
Estimator

Bring Your Own Container - Steps

Code



Containerize



Register with
Amazon ECR




Create
Estimator


- `aws ecr create-repository --repository-name "${algorithm_name}" > /dev/null`
- `fullname="${account}.dkr.ecr.${region}.amazonaws.com/${algorithm_name}:latest"`
- `docker push ${fullname}`

Bring Your Own Container - Steps

Code



Containerize



Register with
Amazon ECR



Create
Estimator

Bring Your Own Container - Steps

Code



Containerize



Register with
Amazon ECR



Create
Estimator

- `byoc_image_uri = '{}.dkr.ecr.{}.{}' / {}'.format(account_id, region, uri_suffix, ecr_repository + tag)`
- `estimator = Estimator (image_name=byoc_image_uri,)`

Summary



Summary

- Tune and evaluate a model
- Model tuning
- Tune a BERT-based text classifier
- Model evaluation
- Evaluate a BERT-based text classifier
- TODO: Script Mode?
- TODO: Bring Your Own Container?

